


CS2102 Database Sys Summary

AY23/24 Sem 1, github.com/gerteck

Topics & Objectives

- **Design:** Entity-Relationship (ER) Model, Functional Dependencies, Normal Forms
- **Implementation:** SQL (Data definition language, Queries, Stored procedures, Triggers)
- **Theory:** Relational Calculus and algebra
- Module covers fundamental concepts and techniques for:
 - Understanding and practice of design & implementation of database applications and management of data with relational db management systems.
 - Design of ER data models to capture data requirements, translate to relational database schema, refine using schema decompositions to avoid anomalies.
 - Use SQL to define relational schemas, write queries.
 - Reason about correctness using concepts of formal query lang (relational calculus & algebra) and apply knowledge to develop database applications.



1. Database Management Sys DBMS

Challenges for Data-Intensive applications

- **Efficiency:** Fast access to information in volumes of data
- **Transactions:** "All or nothing" changes to data
- **Data Integrity:** Parallel access and changes to data
- **Recovery:** Fast and reliable handling of failures (e.g. HD-D/Sys crash, power outage, network disruption)
- **Security:** Fine-grained data access rights

File-based data management to DBMS

- Complex, low level code, Often similar requirements across different programs
- **Problems:** High development effort, Long development times, Higher risk of (critical) errors
- **DBMS:** Set off universal and powerful functionalities for data management, with faster application development, higher stability, less errors.

Core concepts of DBMS

- **ACID Transaction:** Finite sequence of database operations (reads and/or writes), smallest logical unit of work
- **Atomicity:** either all effects of T are reflected in the database or none ("all or nothing")
- **Consistency:** the execution of T guarantees to yield a correct state of the database
- **Isolation:** execution of T is isolated from the effects of concurrent transactions
- **Durability:** after commit of T, its effects are permanent even in case of failures

Concurrent Execution

Concurrent Execution — Common Problems

T _i (B, 500)	T _j (B, 100)
begin	
read(B) <i>100</i>	
B = B + 500 <i>1500</i>	
	begin
	read(B) <i>100</i>
	B = B + 100 <i>1100</i>
	write(B) <i>150</i>
	commit
	write(B) <i>1100</i>
	commit

Final balance B = 1,100 (effect of T_i, overwritten)

→ Lost Update

T _i (B, 500)	T _j (B, 100)
begin	
read(B) <i>100</i>	
B = B + 500 <i>150</i>	
	begin
	read(B) <i>100</i>
	B = B + 100 <i>1100</i>
	write(B) <i>150</i>
	commit
	read(B) <i>150</i>
	begin
	read(B) <i>100</i>
	B = B + 100 <i>1100</i>
	write(B) <i>1100</i>
	commit

Final balance B = 1,600 (when it should be 1,100)

→ Dirty Read

T _i (B, 500)	T _j (B, 100)
begin	
read(B) <i>100</i>	
B = B + 500 <i>150</i>	
	begin
	read(B) <i>100</i>
	B = B + 100 <i>1100</i>
	write(B) <i>1100</i>
	commit
	read(B) <i>100</i>
	begin
	read(B) <i>100</i>
	B = B + 100 <i>1100</i>
	write(B) <i>1100</i>
	commit


Balance B is retrieved twice but the values differ

→ Unrepeatable Read

Require Serializable transaction execution:

- A concurrent execution of a set of transactions is serializable if this execution is equivalent to some serial execution of the same set of transactions
- Two executions are equivalent if they have the same effect on the data
- **DBMS:** Support concurrent executions of transactions to optimize performance, Enforce serializability of concurrent executions to ensure integrity of data

Data Abstraction



Data Independence



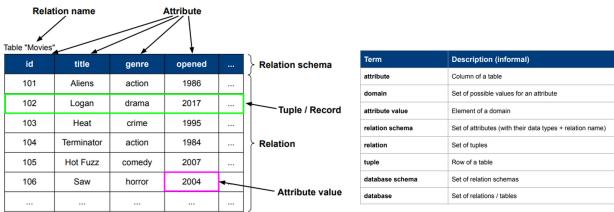
Terminology / Definitions

- **Data Model:** Collection of concepts for describing data
- **Schema:** Description of structure of DB using data model
- **Schema Instance:** Content of a DB at a particular time

Relational Data Model

Data is modelled by relations, and each relation has a definition called a relation schema. This schema specifies attributes (columns) and data constraints (e.g. domain constraints)

- **Relation:** Can be seen as Tables with rows and columns:
 - No. of cols = Degree/Arity, No. of rows = Cardinality
 - Each row is called a tuple/record. It has a component for each attribute of the relation.
 - A relation is thus a set of tuples and an instance of the relation schema, i.e. of a single table.
- **Domain:** Set of atomic values, e.g. integers. All values for an attribute is either in this domain or null.
- **Relational database schema:** Set of relation schemas and their data constraints, i.e. of multiple tables
- **Relational database:** Instance of the schema and is a collection of tables.



Integrity Constraints

Condition that restricts the data that can be stored in a database instance. A legal relation instance is a relation that satisfies all specified ICs.

- **Domain Constraints:** Restrict the attribute values of relations, e.g. only integers allowed

Key Constraints:

- **Superkey:** A superkey is a subset of attributes in a relation that uniquely identifies its tuples.
- **Key:** A key is a superkey which is minimal, i.e. no proper subset of itself is a superkey.
- **Candidate keys:** Set of all possible keys for a relation. One of these keys is selected as the primary key.
- **Primary key:** Chosen candidate key for a relation. Cannot be null (entity integrity constraint). Underlined in relation schema. Prime attribute: Attribute of a primary key (cannot be null)

Foreign Key Constraints:

- **Foreign key:** A foreign key refers to the primary key of a second relation (which can be itself)
- Each foreign key value must be the primary key value in the referenced relation or be null (foreign key constraint)
- Also known as referential integrity constraints.

Term	Description (informal)
(Candidate) key	Minimal set of attributes that uniquely identify a tuple in a relation
primary key	Selected key (in case of multiple candidate keys)
foreign key	Set of attributes that is a key in referenced relation
prime attribute	Attribute of a (candidate) key

- Terminology: DB. vs DBS vs. DBMS

$$\text{DBS} = \text{DBMS} + n \cdot \text{DB} \quad (n > 0)$$

2. SQL: Structured Query Language

- Declarative language: focus on what to compute, not on how to compute
- Contains two parts: Data Definition Language and Data Manipulation Language

Datatypes

type	description
boolean	logical Boolean (true/false)
integer	signed 4-byte integer
float8	double precision floating-point number (8 bytes)
numeric(p, s)	number with p significant digits and s decimal places
char(n)	fixed-length character string
varchar(n)	variable-length character string
text	variable-length character string
date	calendar date (year, month, day)
timestamp	date and time

Integrity Constraints 2

- A **consistent state** of the database is a state which complies with the business rules as defined by the structural constraints and the integrity constraints in the schema.
- If an integrity constraint is violated by an operation or a transaction, the operation or the transaction is aborted and rolled back and its changes are undone, otherwise, it is committed and its changes are effective for all users.
- Five main kinds of integrity constraints in SQL: **NOT NULL, PRIMARY KEY, UNIQUE, FOREIGN KEY, CHECK.**

Primary Key

A primary key is a set of columns that uniquely identifies a record in the table. Each table has at most one primary key. The primary key can be one column or a combination of columns.

```
-- Declare primary key as column constraint
CREATE TABLE customers (
    firstname VARCHAR(64) NOT NULL ,
    lastname VARCHAR(64) NOT NULL,
    email VARCHAR(64) UNIQUE NOT NULL,
    id VARCHAR(16) PRIMARY KEY,
    UNIQUE (firstname, lastname));
```

Composite Primary Key & NOT NULL

A not null constraint guarantees that no value of the column can be set to null. A not null constraint is always declared as a row constraint. When it is explicit, it is declared with the keyword NOT NULL.

```
CREATE TABLE games (
    name VARCHAR(32) ,
    version CHAR(3) ,
    price NUMERIC NOT NULL ,
    PRIMARY KEY (name, version) );
```

Data Insertion Populating Tables

```
INSERT INTO customers VALUES(
    'Carole', 'Yoga', 'cyoga@email.org',
    'Carole89');
```

Deleting Tables

```
DELETE FROM customers
-- DROP deletes content \& table definition
DROP TABLE customers
DROP TABLE IF EXISTS downloads
```

Unique

A unique constraint on a column or a combination of columns guarantees the table cannot contain two records with the same value in the corresponding column or combination of columns.

```
CREATE TABLE customers (
    firstname VARCHAR(64) NOT NULL ,
    lastname VARCHAR(64) NOT NULL,
    email VARCHAR(64) UNIQUE NOT NULL,
    id VARCHAR(16) PRIMARY KEY,
    UNIQUE (firstname, lastname));
```

Foreign Key

- A foreign key constraint enforces **referential integrity**. The values in the columns for which the constraint is declared must exist in the corresponding columns of the referenced table.

- Referenced columns are usually required to be the primary key of the referenced table. Some systems relax this.
- A foreign key is declared using the keyword REFERENCES as a row constraint and the keywords FOREIGN KEY and REFERENCES as a table constraint.

```
CREATE TABLE downloads (
    customerid VARCHAR (16) REFERENCES
        customers (id),
    name VARCHAR (32)
    version CHAR (3),
    FOREIGN KEY (name, version) REFERENCES
        games (name, version) );
```

Check

- Check constraint enforces any other condition that can be expressed in SQL. Declared as row or table constraint.

```
CREATE TABLE games (
    name VARCHAR(32),
    version CHAR(3),
    PRIMARY KEY (name, version)
    price NUMERIC NOT NULL CHECK (price > 0)
    -- or as table constraint:
    CHECK (price > 0) );
```

Update and Delete Propagation

- The annotations ON UPDATE/DELETE with the option CASCADE propagate the update or deletion, when there are chains of foreign key dependencies.

```
CREATE TABLE downloads(
    id VARCHAR (16) REFERENCES customers (id)
        ON UPDATE CASCADE
        ON DELETE CASCADE,
    name VARCHAR(32),
    version CHAR(3),
    PRIMARY KEY (id, name, version),
    FOREIGN KEY (name, version) REFERENCES
        games(name, version)
        ON UPDATE CASCADE
        ON DELETE CASCADE);
```

- Generally a good idea to constraint all columns not to be

null unless there is a good design or tuning reason for not doing so.

- Think carefully about which foreign keys should be subject to cascade.
- Good idea to defer all the constraints that can be deferred. These are checked at the end of a transaction and not immediately after each operation.

Querying Tables

Print Table

- Wildcard '*' to include all attributes

```
SELECT *
FROM customers;
```

View

- We can give a name to a query, called a view. Once created, a view can be queried like a table.
- Creating a view is generally a better option than creating and populating a table, temporary or not.

```
CREATE VIEW sg\customers AS
SELECT c.firstname, c.lastname, c.email,
    c.id
FROM customers c
WHERE country = 'Singapore';

SELECT * FROM sg\customers;
```

3. SQL Queries

Printing one Table

```
SELECT firstname, lastname
FROM customers
WHERE country = 'Singapore';
```

DISTINCT and ORDER BY

- Selecting a subset of columns may result in duplicate row even if original table has a primary key.
- DISTINCT keyword eliminates eventual duplicates, requests results contain distinct rows.
- Both DISTINCT and ORDER BY involve sorting and

conceptually ORDER BY is applied before SELECT DISTINCT.

```
SELECT DISTINCT name, version
FROM downloads
ORDER BY name ASC, version DESC;
```

WHERE

- Returns rows that evaluate to true, filter rows on a Boolean condition
- Uses Boolean operators such as AND, OR and NOT, and various comparison operators such as >, <, >=, <=, <>, IN, LIKE and BETWEEN AND
- Does not return rows that evaluate to unknown/null!
- '-' matches single char
- '%' matches any sequence of zero or more chars

```
SELECT firstname, lastname
FROM customers
WHERE country IN ('Singapore', 'Indonesia')
AND (dob BETWEEN '2001-01-01' AND
      '2000-12-01' OR since >= '2016-12-01 ')
AND lastname LIKE 'B%'
```

- PostgreSQL use "||" for concatenation

```
-- Not to collect GST below 30 cents:
SELECT name || ' ' || version AS game,
       price * 1.07
FROM games
WHERE price * 0.07 < 0.3
```

De Morgan's Laws

```
SELECT name
FROM games
-- all 3 are the same:
WHERE (version = '1.0' or version = '1.1')
WHERE version IN ('1.0', '1.1')
WHERE NOT (version <> '1.0' AND version <>
      '1.1');
```

NULL value

- Every domain has additional value, null. Ambiguous, could be "unknown", "does not exists", or both. In SQL it

is generally (but not always) "unknown".

- With null values, the logic of SQL is a three valued logic with unknown.
- Use `IS (NOT) NULL` for comparison with null
- `COALESCE()` returns the first non-null of its argument.
- `COUNT(*)` counts NULL values.

`COUNT(att) AVG(att) MAX(att) MIN(att)` eliminate null values

Cross Join

- Cross join & Cartesian Product & cross product, represented by comma `CROSS JOIN`,
- Cross join with `WHERE` clause: add condition that FK columns equal to the corresponding PK columns.
- Systematically define table variables (e.g. `games AS g`)

```
SELECT *
FROM customers c, downloads d, games g
WHERE d.id = c.id
AND d.name = g.name
AND d.version = g.version
```

4. Algebraic SQL Queries

Inner Join

- `JOIN` interpreted as `INNER JOIN`
- Inner joins combine records from two tables whenever there are matching values in a field common to both tables.

```
SELECT *
FROM customers c JOIN downloads d ON d.id =
c.id
JOIN games g ON d.name = g.name AND
d.version = g.version;
```

Natural Join

- If we give the same name to columns that are the same, can use natural join. Joins the rows that have the same values for their columns that have the same names. It also prints one of the two equated columns

```
SELECT *
FROM customers c NATURAL JOIN downloads d
NATURAL JOIN games g;
```

Outer Join

- outer join keeps the columns of the rows in the left (left outer join), right (right outer join) or in both (full outer join) tables that do not match anything in the other table according to the join condition and pad the remaining columns with null values.
- Better to avoid outer joins whenever possible as they introduce null values.

`RIGHT (OUTER) JOIN`, `LEFT (OUTER) JOIN`
`FULL (OUTER) JOIN`

```
-- finds customers, never downloaded a game
SELECT c.id FROM customers c
LEFT JOIN downloads d ON c.id = d.id
WHERE d.id IS NULL;
```

Set Operations

- Union, intersect and non-symmetric difference.
- Eliminate duplicates: `UNION`, `INTERSECT`, `EXCEPT`
- Keep duplicates: `UNION ALL`, `INTERSECT ALL`
`EXCEPT ALL`

4. Aggregate SQL Queries

Aggregate Functions

- The values of a column can be aggregated aggregation functions such as `COUNT()`, `SUM()`, `MAX()`, `MIN()`, `AVG()`, `STDDEV()` etc.

```
SELECT COUNT(*)
FROM customers c;
```

```
-- ALL is default and omitted
-- DISTINCT needed
```

```
SELECT COUNT(ALL DISTINCT c.country)
FROM customers c;
```

```
-- Finds min, max, avg and stddev, TRUNC()
-- displays 2 dp
```

```
SELECT MAX(g.price),
MIN(g.price),
TRUNC(AVG(g.price), 2) AS ave,
TRUNC(STDDEV(g.price), 2) AS std
```

```
FROM games g;
```

Group By

- The GROUP BY clause creates groups of records that have the same values for the specified fields before computing the aggregate functions.
- Groups are formed after the rows have been filtered by the WHERE clause
- Recommended (and required by SQL standard) to include attributes projected in the SELECT clause in the GROUP BY clause.
- The order of columns in the GROUP BY clause does not change the meaning of the query.

```
SELECT c.country, COUNT(*)
FROM customers c
WHERE c.dob >= '2000-01-01'
GROUP BY c.country;
```

```
SELECT c.country, EXTRACT( YEAR FROM
c.since) AS regyear, COUNT(*) AS total
FROM customers c, downloads d
WHERE c.id = d.id
GROUP BY c.country, regyear,
ORDERBY regyear, c.country;
```

Having

- Aggregate functions can be used in conditions. However, agg. functions not allowed in WHERE.
- `HAVING` clause to add conditions to be checked after the evaluation of the GROUP BY.
- `HAVING` can only involve aggregate functions, columns listed in the GROUP BY clause and subqueries.

```
SELECT c.country,
FROM customers c
GROUP BY c.country
HAVING COUNT(*) >= 100;
```

5. Nested SQL Queries

Subqueries

- In FROM clause: Must be enclosed in parenthesis, Table alias mandatory, Column aliases optional

- Not recommended, can be written as simple query

```
SELECT cs.lastname, d.name
FROM (SELECT *
      FROM customers c
     WHERE c.country = 'Singapore') AS cs,
          downloads d
WHERE cs.id = d.id;
```

- In WHERE clause, also can be written as simple query.
- Never use a comparison to a subquery without specifying the quantifier ALL or ANY

```
SELECT g1.name, g1.version, g1.price
FROM games g1
WHERE g1.price >= ALL (
    SELECT g2.price
    FROM games g2);
-- or do:
WHERE g1.price = ALL (
    SELECT MAX(g2.price)
    FROM games g2);
-- Note HAVING g.price=MAX(g.price) will
  not work
```

Exists

- **EXISTS** evaluates to true if the subquery has some results.
- Generally correlated. If uncorrelated, then likely either wrong or unnecessary

```
SELECT c.id
FROM customers c
WHERE NOT EXISTS (
    SELECT d.id
    FROM downloads d
    WHERE c.id = d.id);
-- same as
WHERE c.id NOT IN (
    SELECT d.id
    FROM downloads d);
-- same as
WHERE c.id <> ALL (
    SELECT d.id
    FROM downloads d);
```

```
-- Find countries with most customers
SELECT c1.country
FROM customers c1
GROUP BY c1.country
HAVING COUNT(*) >= ALL (
    SELECT COUNT(*)
    FROM customers c2
    GROUP BY c2.country);
```

SQL Conditionals

- CASE expression** • Generic conditional, like if/else statement, Used in SELECT, ORDER BY, etc

```
-- Regular if else
CASE
    WHEN cond1 then res1
    WHEN cond2 then res2
    ...
    WHEN condN then resN
    ELSE res0
END
```

```
-- Switch like statement
CASE expression
    WHEN val1 then res1
    WHEN val2 then res2
    ...
    WHEN valN then resN
    ELSE res0
END
```

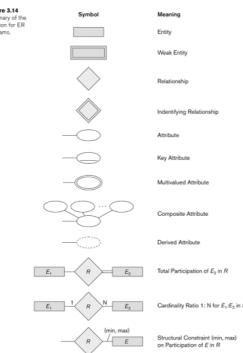
Conceptual evaluation of queries

FROM → WHERE → GROUP BY → HAVING → SELECT → ORDER BY → LIMIT/OFFSET

6. ER Model

Data Modeling Using the Entity-Relationship (ER) Model

NOTATION for ER diagrams

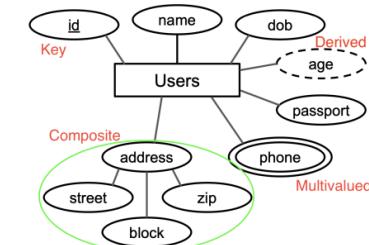


Entity

- Objects that are distinguishable from other objects
- **Entity set:** Collection of entities of the same type
- In ER diagrams, an entity type is displayed in a rectangular box

Attributes

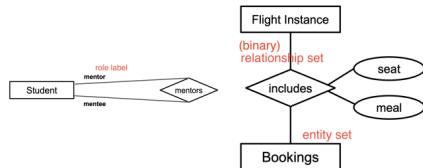
- **Attributes** are properties used to describe an entity
- Each attribute has a value set (or data type) associated with it – e.g. integer
- **Key attribute(s)** uniquely identifies each entity
- **Composite attribute** composed of multiple other attributes
- **Multivalued attribute** may consist of more than one value for a given entity
- **Derived attribute** derived from other attributes



- Attributes are displayed in ovals, multivalued attributes double ovals

Relationship

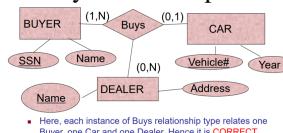
- A **relationship** relates two or more distinct entities with a specific meaning.
- **Degree** of a relationship type is the number of participating entity types. A n-ary relationship set involves n entity roles. Typically binary or ternary.



- **Relationship Set:** Collection of relationships of the same type, can have their own attributes that further describe the relationship
- Most **relationship attributes** are used with M:N relationships: (In 1:N relationships, they can be transferred to the entity type on the N-side of the relationship)
- We represent the relationship type as **Diamond-shaped box**, connected to the participating entity types.
- Relationship type typically readable from left to right and top to bottom.

N-ary relationships


- **Implicit Constraint** In ternary relationship, every instance of relationship must have one instance of each entity.
- Rule extends to n-ary relationships



- **Avoid Ternaries for Easy Modeling:** e.g. “objectifying” the relationship type “Interview” into an entity type ‘Interview’.


Cardinality (Ratio) constraints

- **Upper bound** for entity's participation



Participation / Existence Dependency constraints

- **Lower bound** for entity's participation
- Partial (default): participation not mandatory
- Total: mandatory (at least 1)



Recursive Relationship Type

- A relationship type between the same participating entity type in **distinct roles**.
- In ER diagram, need to display role names to distinguish participations.

Dependency constraints

Weak Entity Types: No key attribute and is identification dependent on another entity.


- Participate in an identifying relationship type with an owner or identifying entity type
- Two kinds of dependencies:

Existence (no weak entity) dependency

Identification (weak entity) dependency

Partial Key

- Set of attributes of weak entity set that uniquely identifies a weak entity, for a given owner entity.



(min,max) notation for relationship constraints

- Read the min,max numbers next to the entity type and looking away from the entity type



7. EER Model

Enhanced ER or Extended ER

IS-A Hierarchies

- “**Is a**” relationship - used to model generalization/specialization of entity sets.
- **Hierarchy:** constraint that every subclass has only one superclass (single inheritance); basically a tree structure.
- **Lattice:** a subclass can be subclass of more than one superclass (multiple inheritance).

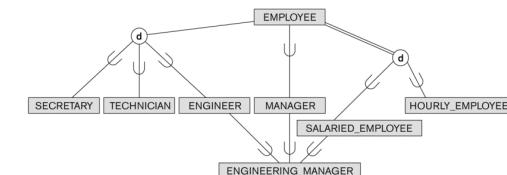


Figure 4.6
A specialization lattice with shared subclass ENGINEERING_MANAGER.

- **Subclass:** subclass member is entity in a **distinct specific role**. Entity inherits all attributes and relationships of superclass.
- **Specialization:** process of defining a set of subclasses of a superclass, based on **distinguishing characteristics**
- **Type of Specialization:** Can be Predicated defined, Attribute defined (Written beside joining line) or User defined.
- **Generalization:** reverse of specialization, based upon some **distinguishing characteristics**.

Constraints on Specialization / Gen

- **Disjointness Constraint:** Subclasses must be disjoint, entity can be member of at most one. (specified by **d** in EER)
- If not disjoint, specialization is **overlapping**, (specified by **o** in EER)
- **Completeness Constraint:** Total specifies every entity must be a member of some subclass. (specified by **double line** in EER)
- **Partial** allows entity not to belong to any subclass. (specified by **single line** in EER)
- Hence, we have four types: Disjoint/Overlapping x Total/Partial. Note generalization usually is total.

8. Relational Mapping

- Preserve info, maintain constraints, minimize null.

1. Map Regular Entity Types

- Create new table (relation R), include all simple attributes.

2. Map Weak Entity Types


- Weak entity type, create table with corresponding FK.

3. Map Binary 1:1 Relationship Type

- Option 1: 2 Foreign Key (2 relations tables) option.
- 2: Merged relation option: Combine relationship set and either entity set into **one** table
- 3: Cross-reference (3 relations, one lookup table).

4. Map Binary 1:N Relationship Type

- For 1:N create table of participating entity (N-side).
- Include FK in table & attributes of the 1:N relation type.



5. Map Binary N:M Relationship Type

- Represent relationship set with a new table. Include as FK the PK of relations, combination will form PK of table.

6. Map Multi-valued Attributes

- Create new table / relation for each multivalued attribute.

7. Map N-ary relationship type

- For each n-ary relationship type R, where n > 2, create new table. Include FK of relations, and attributes.

Summary:

Table 9.1 Correspondence between ER and Relational Models

ER MODEL	RELATIONAL MODEL
Entity type	Entity relation
1:1 or 1:N relationship type	Foreign key (or <i>relationship</i> relation)
M:N relationship type	<i>Relationship</i> relation and two foreign keys
<i>n</i> -ary relationship type	<i>Relationship</i> relation and <i>n</i> foreign keys
Simple attribute	Attribute
Composite attribute	Set of simple component attributes
Multivalued attribute	Relation and foreign key
Value set	Domain
Key attribute	Primary (or secondary) key

8. Map Specialization or Generalization

- Option 1: Multiple tables (relations) - Superclass and subclasses
- 2: Table each for each subclass relations, inherit superclass attributes
- 3: Single table with one type attribute
- 4: Single table with multiple type attributes
- For specialisation hierarchies with one superclass and n subclasses, possibility of mapping from 1 relation to (n+1) relations. Design highly subjective, try to maintain appropriate attributes to determine subclass identity.

9. Relational Calculus & Algebra

- Both are formal query languages. A query is composed of a collection of operators called relational operators.
- Relational Calculus** (declarative language: what is to be done rather than how to do it). Order not specified, concerned with result we have to obtain.
- Relational Algebra**: (procedural language) Order specified in which operations have to be performed.
- Relational algebra basis of implementation of relational DBMS.
- SQL queries translated into execution plans, orchestrate physical implementation of relational algebra operators.

Predicate Logic

- Predicate / first order logic:** formulae built from:
- predicates** (lowercase), **operators** ($=$, etc.) **constants** (lower case), **variables** (upper case, quantified or free), **connectives** (\rightarrow), and **quantifiers** (\forall and \exists).
- Existential quantifier**, \exists : disjunction:

$$\exists X, F[X] \equiv F[a] \vee F[b] \vee \dots$$
- The universal quantifier**, *forall*: conjunction.

$$\forall X, F[X] \equiv F[a] \wedge F[b] \wedge \dots$$
- De Morgan laws** applies to quantifiers.

$$\neg(\exists X, F[X]) \equiv \forall X, \neg F[X]$$

$$\neg(\forall X, F[X]) \equiv \exists X, \neg F[X]$$

Relational Calculus (Tuple)

- Concerned with **result we have to obtain**.
- In **Tuple Relation Calculus (TRC)** variables values of rows of tables (tuples), is the theoretical basis of SQL.
- Material Implication:** $P \rightarrow Q \equiv \neg P \vee Q$
- Relational Calculus denoted as:**

$$\{t | P(t)\}$$

- t**: set of tuples, **P**: condition which is true for the given set of tuples.
- E.g. $\{T | \exists T_1 (T_1 \in \text{department} \wedge T_1.\text{faculty} = \text{'School of Computing'} \wedge T.\text{department} = T_1.\text{department})\}$

Relational Algebra

- Order is specified in which operations have to be performed.

Unary Operators

Selection, σ_c

- For each tuple $T \in R, T \in \sigma_c(R)$, means selection condition c evaluates to true for tuple t .
- E.g. Find the customers from Singapore.

$$\sigma_{c.\text{country} = \text{'Singapore'}}(\rho(\text{customers}, c))$$

- Equivalent to:

```
SELECT * FROM customers c
WHERE c.country = 'Singapore'
```

- Condition** is boolean expression of form:

expression	example
attribute op constant	$\sigma_{\text{start}=2020}(\text{Projects})$
<i>attr₁</i> op <i>attr₂</i>	$\sigma_{\text{start}=\text{end}}(\text{Projects})$
<i>expr₁</i> \wedge <i>expr₂</i>	$\sigma_{\text{start}=2020 \wedge \text{end}=2021}(\text{Projects})$
<i>expr₁</i> \vee <i>expr₂</i>	$\sigma_{\text{start}=2020 \vee \text{end}=2021}(\text{Projects})$
$\neg \text{expr}$	$\sigma_{\neg(\text{start}=2020)}(\text{Projects})$
(<i>expr</i>)	-

- op** $\in \{=, <, <, \leq, \geq, >\}$
- Precedence: (), **op**, \neg , \wedge , \vee
- null** comparison is **unknown**, arithmetic with **null** is **null**

Projection π_l

- Projects columns of a table specified in **list l**.
- (**SELECT xx, yy FROM games**)
- Order of attribute in l matters**.
- Duplicates removed.
- Examples:**

$$\pi_{g.name, g.version, g.price}(\rho(games, g))$$

Teams		
en	pn	hours
Sarah	BigAI	10
Sam	BigAI	5
Sam	BigAI	3

$\pi_{pn, en}(\text{Teams})$	
pn	en
BigAI	Sarah
BigAI	Sam

Renaming, ρ_l

- Can change name of relation, of attributes, or both.
- Change name of relation: $\rho(R_1, R_2)$
- Change attribute names: $\rho(R_1, R_1(a_1 \rightarrow b_1, a_2 \rightarrow b_2))$

Set Operations

- Set operations include \cup, \cap, \times , set difference (\setminus)
- Intersection able to express with union and set difference:

$$R \cap S = (R \cup S) - ((R - S) \cup (S - R))$$

- Union Compatability:** two relations must be union compatible. Have same number of attributes, corresponding attributes have same or compatible domains.
- (i.e. relations must have same columns).
- Cross Product:** (Cartesian Product) Forms all possible pairs of tuples from two relations.

$$R_1 \times R_2$$

Join Operations

- Combines \times, σ_c, π_l into a single op.
- Simple relational algebra expressions

Inner Joins

- Eliminates tuples that do not satisfy matching criteria (i.e. selection)
- Is a selection from cross product

$$R \bowtie_C S = \sigma_C (R \times S)$$
- Example:

$$\rho(\text{customers}, c) \bowtie_{d.id=c.id} \rho(\text{downloads}, d)$$

Relational Calculus & Algebra

- 4 Steps to construct calculus and algebra queries:**
 - Construct SQL query you are familiar with (difficult)
 - From query, map the tables that you need (yellow)
 - From query, map the conditional statements (blue)
 - From query, map the columns you need to print (green)

Q1 (A)	Q: Find the different departments in School of Computing
	$(T_1 \exists T_1.T_1 \in \text{department} \wedge T_1.\text{faculty} = \text{'School of Computing'}) \wedge T_1.\text{department} = T_1.\text{department}$
Q2 (A)	Q: Find the different departments in School of Computing.
	$\rho(\text{department}, d) \bowtie_{d.faculty = \text{'School of Computing'}} \rho(\text{department}, d)$
Equivalent Query	Equivalent Query
SELECT DISTINCT d.department FROM department d WHERE d.faculty = 'School of Computing';	SELECT DISTINCT d.department FROM department d WHERE d.faculty = 'School of Computing';

10. Programming with SQL

Writing Database Applications

- **Interactive SQL:** Directly writing SQL statements to an interface. (e.g. PostgreSQL's psql cli, pgAdmin).
 - **Non-interactive SQL:** SQL statements included in application written in host language.
 - 2 Main alternatives: **Statement Level Interface (SLI), and Call Level Interface (CLI).**
 - Crudely, SLI = CLI in disguise, as SLI preprocessor generates CLI code.

Statement Level Interface (SLI)

- Code is mix of host language statements and SQL statements (e.g. embedded SQL, dynamic SQL).
 - **Basic process for SLI:** Write code that mixes host language with SQL, preprocess code using a preprocessor, compile code into exe program.

Statement Level Interface (SLI)

SLI — Common Steps

- **Declaration**
 - Declare variables
 - Can be used by host language & SQL (other variables only usable by host language)
 - **Connection**
 - Connect to database with credentials
 - **Execution**
 - Prepare queries (in case of dynamic SQL)
 - Execute queries (might require cursors)
 - Operate on result
 - **Deallocation**
 - Release all resources
 - If needed: commit any changes to db

```

Source file: staticquery.pgc
[1] main()
{
    EXEC SQL DECAL DECLARE SECTION;
    cur V_HSTL1 CURSOR FOR SELECT * FROM V_HSTL1 WHERE HSTL1='localhost';
    EXEC SQL EXECUTE IMMEDIATE cur;
    EXEC SQL CLOSE cur;
    EXEC SQL DEALLOCATE SECTION;
}

// Declare CONNECT to target user using -password-
EXEC SQL CONNECT TO target_user user USING '-password';


// Declare cursor cur using HSTL1
EXEC SQL DECLARE cur CURSOR FOR
SELECT * FROM V_HSTL1 WHERE HSTL1='localhost' LEFT JOIN
    // Open cursor
EXEC SQL OPEN cur;
EXEC SQL WHENEVER NOT FOUND DO BROWSE;

// Fetch through cursor and display results
fetch|||;
    EXEC SQL PULL cur INTO :i_hstl1, :i_hstl2, :i_hstl3;
}
    // Close cursor, commit, disconnect
    EXEC SQL CLOSE cur;
    EXEC SQL COMMIT;
    EXEC SQL DISCONNECT;
}

return 0;
}

```

SII — Preprocessing, Compiling, Running Code



SII — Dynamic SII

- **Dynamic SQL**
 - SQL query is generated at runtime
 - Example on the right: number of riders are specified as command line parameter

```
[base] wduffy@script: /Users/wduffy/Desktop/100-days [base] wduffy@script: /Users/wduffy/Desktop/100-days [base] ./dynamicus-q15  
[base] b102_ 1: name: Tiers Benoit  
[base] b102_ 2: name: Tiers Benoit  
[base] b102_ 3: name: Christophe Laporte  
[base] b102_ 4: name: Søren Kragh  
[base] b102_ 5: name: Christophe Laporte  
[base] b102_ 6: name: Søren Kragh  
[base] b102_ 7: name: Dylan Van Barel  
[base] b102_ 8: name: Natan Von Hoeydorp  
[base] b102_ 9: name: Søren Kragh  
[base] b102_ 10: name: Dylan Van Barel  
[base] b102_ 11: name: Natan Von Hoeydorp  
[base] b102_ 12: name: Mikkil Bjørk  
[base] b102_ 13: name: Søren Kragh  
[base] b102_ 14: name: Søren Kragh  
[base] b102_ 15: name: Vesparé Stale Lækken  
[base] b102_ 16: name: Rafal Majka  
[base] b102_ 17: name: Søren Kragh  
[base] b102_ 18: name: Matteo Trentin
```

Call Level Interface (CLI)

- Application completely written in host language, while **SQL statements** are **strings** passed as **arguments** to host language procedures or libraries
 - E.g. ODBC (Open DataBase Connectivity), JDBC (Java DB Connectivity), psycopg library for Python - PostgreSQL.

CLI — Static SQL Example

```

Declaration           import psycopg2 # PostgreSQL library (here psycopg for Python)

Connection          # Connect to database
connection = psycopg.connect("host=local dbname=timescale123 user=postgres password=demo")

# Create cursor
cursor = connection.cursor()

# Execute query by executing query (string parameter passed in execute() method)
cursor.execute("SELECT bib_id, name FROM books LIMIT 10")

Execution          # Loop over all results until no next tuple is returned
while True:
    row = cursor.fetchone()
    if not row:
        break
    print(f"==> {row[0]}, {row[1]}")

cursor.close()
connection.close()

Deallocation        >>> bib_1, name: Jonas Vingeland
>>> bib_2, name: Kristoffer Kjellberg
>>> bib_3, name: Wilco Kelderman

```

CLI — Dynamic SQL Example

```

Declaration import psycopg2 # Most language library (here psycopg for Python)
Connection {
    # A connect to database
    connection = psycopg.connect("host=localhost dbname=sql2017_tdfb user=ndem password=demo")
}

Execution {
    # A create cursor
    cursor = connection.cursor()
    cursor.execute("SELECT * FROM TBL_Note", [1001])
    rows = cursor.fetchall()
    # A loop over all results until no next tuple is returned
    while True:
        row = cursor.fetchone()
        if row is None:
            break
        print("%s %s %s" % (row[0], row[1], row[2]))
}

Deallocution {
    # cursor.close()
    # connection.close()
}

```

SQL Injection Attack

- Class of cyber attacks on dynamic SQL, goal is to execute unintended (malicious) SQL statements.
 - **Typical cause:** dynamic queries are generated by merging / concatenating strings.
 - **Common attack point:** Omnipresent form fields in web interfaces. Entered values define some SQL statement.
 - **Key Points:** Don't manually merge values to a query, don't use % or + operator to merge values, use provided methods.

11. SQL Functions and Procedures


- Tasks requiring multiple DB operations common, involve any combination of reads and writes.
- E.g. update user password: check user exists → check new password differs from old → if ok, update password (3 separate requests/accesses to DB)
- **Problems:** Application and DB may run on different machines, poor performance or DB becomes bottleneck.
- Different DB operations only loosely connected, difficult to ensure “all or nothing” behavior.
- **Approach:** Move (some) application logic into DB, group DB operations that form task together, treat task as single DB operation.

Stored Functions and Procedures

- **Collection of SQL statements and procedural logic,** precompiled and reusable code, allows execute multiple database operations as a single unit.
- **Procedural Logic:** Relevant for application logic that requires assignments, conditionals or loops, and queries that cannot be expressed using basic SQL.
- **ISO standard:** SQL/PSM (Persistent Stored Modules). Different DBMS have their own flavor.
- **Advantages:** better performance, code reuse, ease of maintenance, added security.
- **Disadvantages:** testing & debugging more challenging, limited portability / vendor lock in, no simple versioning of code, not the most intuitive language.

Stored Functions

Stored Functions: Basic Syntax



Function

```

CREATE OR REPLACE
    FUNCTION <name>
(<param> <type>, ...)
RETURNS <type> AS $$
    <code>
$$ LANGUAGE <sql | plpgsql>;

```

Procedure

```

CREATE OR REPLACE
    PROCEDURE <name>
(<param> <type>, ...) AS
    $$ 
    <code>
$$ LANGUAGE <sql | plpgsql>;

```

CREATE

OR REPLACE

helps to re-declare function/procedure if already previously defined

- Code is enclosed within `$$`
- Call a function:
`SELECT * FROM swap(2, 3);`
- Call a procedure:
`CALL transfer('Alice', 'Bob', 100);`

Return and Type

Return	Type
Single tuple from table	<code><table_name></code>
Set of tuples from table	<code>SET OF <table_name></code>
Single new tuple	<code>RECORD</code>
Set of new tuples	<code>SET OF RECORD</code> or <code>TABLE(c VARCHAR, x INT)</code>
No return value	<code>VOID</code> , or use <code>PROCEDURE</code> instead of <code>FUNCTION</code>
Trigger	<code>TRIGGER</code>

Control Structures

- **Conditionals:**
 - 4 types of `IF` expressions
 - `IF ... THEN ... END IF`
 - `IF ... THEN ... ELSE ... END IF`
 - `IF ... THEN ... ELSIF ... THEN ... ELSE ... END IF`
 - 2 types of `CASE` expressions
 - `CASE ... WHEN ... THEN ... ELSE ... END CASE`
 - `CASE WHEN ... THEN ... ELSE ... END CASE`
- **Simple Loops**
 - `LOOP ... END LOOP` (typically requires `EXIT WHEN` to jump out of loop)
 - `WHILE ... LOOP ... END LOOP`
 - `FOR ... IN ... LOOP ... END LOOP`

Variables

- `DECLARE [var, type]` (1 or more when `DECLARE` keyword is present)
- `<var> := <expr>`
- `BEGIN ... END`
- For plpgsql, code in the BEGIN-END block is in a transaction

Examples

Note: `INOUT` specifies that the param is both an input and output param

Function

```

CREATE OR REPLACE
    FUNCTION swap(INOUT val1
                  INT, INOUT val2
                  INT)
RETURNS RECORD AS $$ 
DECLARE
    temp INT;
BEGIN
    temp := val1;
    val1 := val2;
    val2 := temp;
END;
$$ LANGUAGE plpgsql;

```

Procedure

```

CREATE OR REPLACE
    PROCEDURE transfer(
src TEXT, dst TEXT,
amt NUMERIC
) AS $$ 
UPDATE Accounts
SET balance =
balance - amt
WHERE name = src;
UPDATE Accounts
SET balance =
balance + amt
WHERE name = dst;
$$ LANGUAGE sql;

```

Stored Procedures

- **Stored procedures**
 - Essentially the same syntax as for functions
 - Most obvious difference: not `RETURNS` clause
 - Invoked using `CALL` command

```

CREATE PROCEDURE add_bonus_proc(sid INT, amount INT)
AS
$$
    UPDATE students
    SET points = points + amount
    WHERE id = sid;
$$
LANGUAGE sql;

```

- **Functions** must return something (but can be `VOID`).
- **Procedures** do not have to return anything, but can (using `INOUT` and `OUT` params).

Stored Functions vs. Procedures

- Procedures can **commit or roll back transactions** during execution, functions can not.
 - Unlike functions, **procedures cannot be involved in DML commands** (select, insert, update, delete)
 - Procedures are invoked in isolation using **CALL** (functions always invoked in **SELECT** statements).
 - **Best practice:** return value(s): create function, no return value: create procedure.

Assignments

- Basic Assignment with `:=`, e.g. `age := 29;`
 - Assignment of query result to declared variable(s):
`SELECT ... INTO ...`

Errors & Messages


- RAISE
 - 6 different raise levels in PostgreSQL.

RAISE DEBUG	
RAISE LOG	<ul style="list-style-type: none">Generate messages of different priority levels
RAISE INFO	<ul style="list-style-type: none">Whether messages of a particular priority are reported to the client, depends on the PostgreSQL configuration
RAISE NOTICE	
RAISE WARNING	
RAISE EXCEPTION	<ul style="list-style-type: none">Raises an errorTypically aborts current transaction

Cursors

- **Purpose:** Encapsulates query, allows access to each individual row return by a SELECT. (instead of executing whole query to get complete result at once)
 - **Avoids memory overrun** when the query result is very large.
 - **Advantage:** Flexible ‘navigation’ through query results in **different directions**. `FETCH` to move row and read data. `MOVE` only to move to row (no read)

- General workflow



A cursor is associated with a **SELECT** statement in the declaration block

Cursor Directions

NEXT	Fetch the next row (default)
PRIOR	Fetch the prior row
FIRST	Fetch the first row of the query (same as ABSOLUTE 1)
LAST	Fetch the last row of the query (same as ABSOLUTE -1)
ABSOLUTE <i>n</i>	<ul style="list-style-type: none"> Fetch the <i>n</i>-th row of the query, if <i>n</i> >= 0 Fetch abs(<i>n</i>)-th row from the end, if <i>n</i> < 0. ABSOLUTE 0 positions before the first row
RELATIVE <i>n</i>	<ul style="list-style-type: none"> Fetch the <i>n</i>-th succeeding row, if <i>n</i> >= 0 Fetch the abs(<i>n</i>)-th prior row, if <i>n</i> < 0 Position before first row or after last row if <i>n</i> is out of range RELATIVE 0 re-fetches the current row, if any
FORWARD	Fetch the next row (same as NEXT)
BACKWARD	Fetch the prior row (same as PRIOR).

Question

Given the table "Scores" from before, write a function to perform the following task:

1. Sort the students in "Scores" in descending order of their **Mark** (*break ties arbitrarily*)
 2. For each student, compute the *difference* between his/her **Mark** and the **Mark** of the previous student
 - If there is no previous student, use **NULL**.

Solution

```

CREATE OR REPLACE FUNCTION score_gap()
RETURNS TABLE(name TEXT, mark INT, gap INT) AS $$

DECLARE
    curs CURSOR FOR (SELECT * FROM Scores ORDER BY Mark DESC);
    r RECORD;
    prev INT;
BEGIN
    prev := -1; OPEN curs;
    LOOP
        FETCH curs INTO r;
        EXIT WHEN NOT FOUND;
        name := r.Name; mark := r.Mark;
        IF prev >= 0 THEN gap := prev - mark;
        ELSE gap := NULL;
        END IF;
        RETURN NEXT r; — Insert into output
        prev := r.mark;
    END LOOP;
    CLOSE curs;
END;
$$ LANGUAGE plpgsql;

```

Explanation

1. Declare a **cursor** associated with a **SELECT statement**
 - it's a **RECORD** to store previous rows
 2. **Open** the cursor which executes the **SQL statement** and let the cursor point to the **beginning** of the result.
 3. **Fetch** a tuple from the cursor by reading the **next tuple** from cursor and assign into the variable
 4. If the **FETCH operation** did not get any tuple, the loop **terminates**
 - Otherwise, perform the **main operation** and insert into output
 5. **Close** the cursor to release the resources allocated.


Examples

Cursors — Example (beyond NEXT)

```

CREATE OR REPLACE FUNCTION median_points()
RETURNS NUMERIC AS
$$
DECLARE
    c CURSOR FOR (SELECT * FROM students ORDER BY points DESC);
    s1 RECORD; s2 RECORD; num_students INT;
BEGIN
    OPEN c;
    SELECT COUNT(*) INTO num_students FROM students;
    IF num_students=2 THEN
        FETCH ABSOLUTE (num_students+1)/2 FROM c INTO s1;
        RETURN s1.points;
    ELSE
        FETCH ABSOLUTE num_students/2 FROM c INTO s1;
        FETCH NEXT FROM c INTO s2;
        RETURN (s1.points+s2.points)/2;
    END IF;
    CLOSE c;
END;
$$
LANGUAGE plpgsql;

```



Dynamic Cursors — Example

```

CREATE OR REPLACE FUNCTION median_points(IN has_graduated BOOLEAN,
                                         RETURNS NUMERIC AS
$-
DECLARE
    c CURSOR (grad BOOLEAN) FOR (SELECT * FROM students
                                    WHERE graduated = grad
                                    ORDER BY points DESC);
    s1 RECORD; s2 RECORD; num_students INT;
BEGIN
    OPEN c FOR (SELECT * FROM students
                WHERE graduated = has_graduated);
    SELECT COUNT(*) INTO num_students
    FROM students WHERE graduated = has_graduated;
    IF num_students%2 = 1 THEN
        FETCH ABSOLUTE (num_students+1)/2 FROM c INTO s1;
        RETURN s1.points;
    ELSE
        FETCH ABSOLUTE num_students/2 FROM c INTO s1;
        FETCH NEXT FROM c INTO s2;
        RETURN (s1.points+s2.points)/2;
    END IF;
    CLOSE c;
END;
$-
LANGUAGE plpgsql;

```

