


# CS2102 Database Sys Summary

AY23/24 Sem 1, github.com/gerreck

## Topics & Objectives

- **Design:** Entity-Relationship (ER) Model, Functional Dependencies, Normal Forms
- **Implementation:** SQL (Data definition language, Queries, Stored procedures, Triggers)
- **Theory:** Relational Calculus and algebra
- Module covers fundamental concepts and techniques for:
  - Understanding and practice of design & implementation of database applications and management of data with relational db management systems.
  - Design of ER data models to capture data requirements, translate to relational database schema, refine using schema decompositions to avoid anomalies.
  - Use SQL to define relational schemas, write queries.
  - Reason about correctness using concepts of formal query lang (relational calculus & algebra) and apply knowledge to develop database applications.



## 1. Database Management Sys DBMS

### Challenges for Data-Intensive applications

- **Efficiency:** Fast access to information in volumes of data
- **Transactions:** "All or nothing" changes to data
- **Data Integrity:** Parallel access and changes to data
- **Recovery:** Fast and reliable handling of failures (e.g. HD-D/Sys crash, power outage, network disruption)
- **Security:** Fine-grained data access rights

## File-based data management to DBMS

- Complex, low level code, Often similar requirements across different programs
- **Problems:** High development effort, Long development times, Higher risk of (critical) errors
- **DBMS:** Set off universal and powerful functionalities for data management, with faster application development, higher stability, less errors.

## Core concepts of DBMS

- **ACID Transaction:** Finite sequence of database operations (reads and/or writes), smallest logical unit of work
- **Atomicity:** either all effects of T are reflected in the database or none ("all or nothing")
- **Consistency:** the execution of T guarantees to yield a correct state of the database
- **Isolation:** execution of T is isolated from the effects of concurrent transactions
- **Durability:** after commit of T, its effects are permanent even in case of failures

## Concurrent Execution


### Concurrent Execution — Common Problems

| $T_1(B, 500)$  | $T_2(B, 100)$  | $T_3(B, 500)$   | $T_4(B, 100)$      |
|--|--|---|--------------------|
| begin  | begin  | begin   | begin              |
| read(B) 1000   | read(B) 1000   | read(B) 1000  | read(B) 1000       |
| $B = B + 500$ 1500   | $B = B + 500$ 1500                                     | $B = B + 500$ 1500                                      | $B = B + 500$ 1500 |
| write(B) 1500  | write(B) 1500  | write(B) 1500   | write(B) 1500      |
| commit   | commit   | commit  | commit             |
| Final balance $B = 1,100$<br>(effect of $T_1$ overwritten) | Final balance $B = 1,600$<br>(when it should be 1,100) | Balance $B$ is retrieved twice<br>but the values differ |                    |
| → Lost Update  | → Dirty Read   | → Unrepeatable Read                                     |                    |


Require Serializable transaction execution:

- A concurrent execution of a set of transactions is serializable if this execution is equivalent to some serial execution of the same set of transactions
- Two executions are equivalent if they have the same effect on the data
- **DBMS:** Support concurrent executions of transactions to optimize performance, Enforce serializability of concurrent executions to ensure integrity of data

## Data Abstraction



## Data Independence



## Terminology / Definitions

- **Data Model:** Collection of concepts for describing data
- **Schema:** Description of structure of DB using data model
- **Schema Instance:** Content of a DB at a particular time

## Relational Data Model

Data is modelled by relations, and each relation has a definition called a relation schema. This schema specifies attributes (columns) and data constraints (e.g. domain constraints)

- **Relation:** Can be seen as Tables with rows and columns:
  - No. of cols = Degree/Arity, No. of rows = Cardinality
  - Each row is called a tuple/record. It has a component for each attribute of the relation.
  - A relation is thus a set of tuples and an instance of the relation schema, i.e. of a single table.
- **Domain:** Set of atomic values, e.g. integers. All values for an attribute is either in this domain or null.
- **Relational database schema:** Set of relation schemas and their data constraints, i.e. of multiple tables
- **Relational database:** Instance of the schema and is a collection of tables.

| Table "Movies" |            |        |        |     |                 |
|----------------|------------|--------|--------|-----|-----------------|
| id             | title      | genre  | opened | ... |                 |
| 101            | Aliens     | action | 1986   | ... | Relation schema |
| 102            | Logan      | drama  | 2017   | ... |                 |
| 103            | Heat       | crime  | 1995   | ... | Tuple / Record  |
| 104            | Terminator | action | 1984   | ... |                 |
| 105            | Hot Fuzz   | comedy | 2007   | ... | Relation        |
| 106            | Saw        | horror | 2004   | ... |                 |
| ...            | ...        | ...    | ...    | ... | Attribute value |

| Term            | Description (informal)                                    |
|-----------------|---|
| attribute       | Column of a table   |
| domain          | Set of possible values for an attribute                   |
| attribute value | Element of a domain                                       |
| relation schema | Set of attributes (with their data types + relation name) |
| relation        | Set of tuples   |
| tuple           | Row of a table  |
| database schema | Set of relation schemas                                   |
| database        | Set of relations / tables                                 |

## Integrity Constraints

Condition that restricts the data that can be stored in a database instance. A legal relation instance is a relation that satisfies all specified ICs.

- **Domain Constraints:** Restrict the attribute values of relations, e.g. only integers allowed
- **Key Constraints:**
  - **Superkey:** A superkey is a subset of attributes in a relation that unique identifies its tuples.
  - **Key:** A key is a superkey which is minimal, i.e. no proper subset of itself is a superkey.
  - **Candidate keys:** Set of all possible keys for a relation. One of these keys is selected as the primary key.
  - **Primary key:** Chosen candidate key for a relation, Cannot be null (entity integrity constraint), Underlined in relation schema. Prime attribute: Attribute of a primary key (cannot be null)
- **Foreign Key Constraints:**
  - **Foreign key:** A foreign key refers to the primary key of a second relation (which can be itself)
  - Each foreign key value must be the primary key value in the referenced relation or be null (foreign key constraint)
  - Also known as referential integrity constraints.

| Term            | Description (informal)   |
|-----------------|--|
| (candidate) key | Minimal set of attributes that uniquely identify a tuple in a relation |
| primary key     | Selected key (in case of multiple candidate keys)                      |
| foreign key     | Set of attributes that is a key in referenced relation                 |
| prime attribute | Attribute of a (candidate) key   |

- Terminology: DB. vs DBS vs. DBMS

$$DBS = DBMS + n \cdot DB \quad (n > 0)$$