

CS3223 Database Systems

Implementation

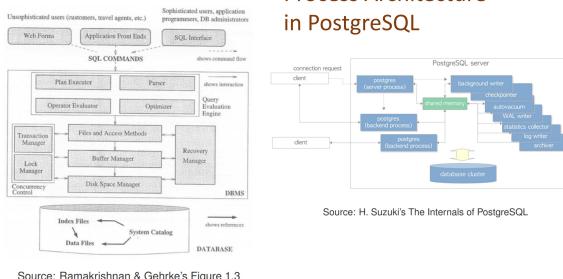
AY23/24 Sem 2, github.com/gerteck

Introduction

Course Details

- Prerequisite knowledges: CS2040S, CS2102, CS2106 background (helpful).
- Reference Textbook: Raghu & Johannes Database M. Systems, 2002. Encouraged to read ahead based on schedule before the lecture.
- Course covers data structures, algorithms, different components making up database systems.

Architecture of DBMS



Source: Ramakrishnan & Gehrke's Figure 1.3

- OLTP:** Online Transaction Processing is a type of data processing that consists of executing a number of transactions occurring concurrently—online banking, shopping, order entry, or sending text messages, for example.
- OLAP:** Online Analytical Processing.
- Focusing on centralized database running on a single server.

1. Data Storage

References: R&G Chapt 8. (Storage & Indexing Overview), Chapt 9. (Storing Data: Disks and Files).

A DBMS stores

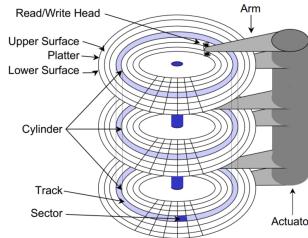
- Relations (Actual tables)
- System catalog (aka data dictionary) storing metadata about relations. (Relation schemas, structure of relations, constraints, triggers. View definitions, Indexes - derived info to speed up access to relations, Statistical information about relations for use by query optimizer.)
- Log files: Information maintained for data recovery.

DBMS Storage

Memory Hierarchy: Primary (registers, RAM), secondary (HDD, SSD), tertiary memory with capacity / cost / access speed / volatility tradeoffs.

- DBMS stores data on non-volatile disk for persistence.
- DBMS processes data in main memory (RAM).
- Disk access operations (I/O). Read: transfer data from disk to RAM. Write: transfer data from RAM to disk.
- Make use of index to speed up access, so that don't have to retrieve all the data when you run a query. Retrieve index and read only the block that contains specified data. Minimize I/O cost.

Magnetic Hard-Disk Drive HDD



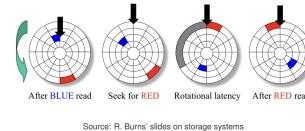
Source: R. Burns' slides on storage systems

- Cylinder, Track, Sector: Units of the HDD storage system. To read from different tracks, need to move the mechanical HDD arm.

Disk Access Time:

- command processing time: interpreting access command by disk controller.
- seek time: moving arms to position disk head on track.
- rotational delay: waiting for block to rotate under head.
- transfer time: actually moving data to/from disk surface.
- access time:** = seek time + rotational delay + transfer time. (CPT considered negligible).

Disk Access Time Components



- Seek time
 - avg. seek time: 5-6 ms
- Rotational delay (or rotational latency)
 - Depends on rotation speed - measured in rotations per minute (RPM)
 - Average rotational delay = time for $\frac{1}{2}$ revolution
 - Example: For 10000 RPM, avg. rotational delay = $0.5 \times (60 / 10000) = 3$ ms
- Transfer time
 - n = number of requested sectors on same track
 - transfer time = $n \times \text{time for one revolution}$
 - avg. sector transfer time: 100-200 μ s
 - Sequential vs random I/O

Concept of Sequential vs random I/O.

- Sequential:** Both sector on same track.
- Random:** Sectors on different track, require seeking (moving arm).
- Given a set of data, we hope to store the data contiguously, on the same track. (Minimize incurring random I/O). If data is too large, store on same track, but different surface (aka same cylinder).
- Complexity hidden to OS by disk controller. Shown as a sequence of memory locations.

Solid-State Drive: SSD

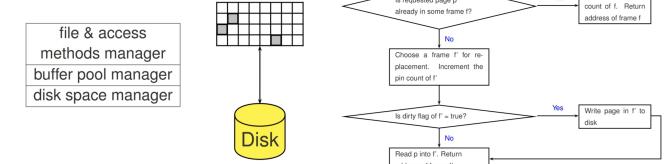
- Build with NAND flash memory without any mechanical moving parts. Lower power consumption.
- Random I/O:** 100x faster than HDD. (no moving parts)
- Sequential I/O:** slightly faster than HDD (2x)
- Disadvantages:** update to a page requires erasure of multiple pages (5ms) before overwriting page. Limited number of times a page can be erased ($10^5 - 10^6$)

Storage Manager Components

- Data is stored, retrieved in units called **disk blocks (or pages)**. (Each block = sequence of one or more contiguous sectors).
- Files & access methods layer (aka file layer)** - deals with organization and retrieval of data.
- Buffer Manager** - controls reading/writing of disk pages.

- Disk Space Manager** - keeps track of pages used by file layer.

Storage Manager Buffer Manager BM: Handling request for page p Components



Buffer Manager

- Buffer pool:** Main memory allocated for DBMS.
- Buffer pool is partitioned into block-sized pages called **frames**.
- Clients of buffer pool can request for disk page to be fetched into buffer pool, release a disk page in buffer pool.
- A page in the buffer is **dirty** if it has been modified & not updated on disk.
- Two variables** maintained for each frame in buffer pool:
 - pin count:** number of clients using page (initialized 0)
 - dirty flag:** whether page is dirty (initialized false)
- Free list: Keeps track of frames that are free / empty.
- Pin count:**
 - Incrementing pin count is **pinning** the requested page in its frame.
 - Decrementing is **unpinning** the page.

- Unpinning a page, dirty flag should be updated to true if page is dirty.
- A page in buffer can be replaced only when pin count is 0.
- Before replacing buffer page, needs to be written back to disk if its dirty flag is true.

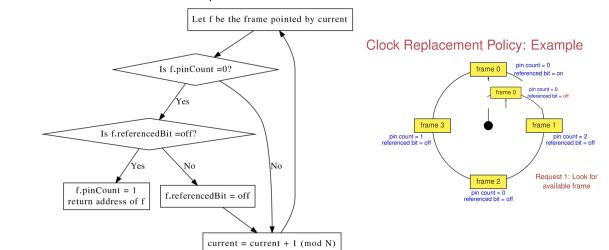
- Buffer manager coordinates with transaction manager to ensure data correctness and recoverability.

Replacement Policies

- Replacement policy: Deciding which unpinned page to replace. (some examples:)
- Random, FIFO, Most Recently Used (MRU), Least Recently Used (LRU): (Use queue of pointers to frames with pin count = 0), most common, makes use of temporal locality.
- Clock:** cheaper popular variant of LRU
 - current** variable: points to some buffer frame.
 - Each frame has a **referenced bit**, turns on when its pin count turns 0.
 - Replace a page that has referenced bit off & pin count = 0.

Clock Replacement Policy

N = number of frames in buffer pool



CS3223: Sem 2, 2023/24

Files

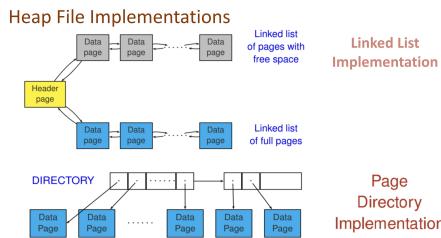
File Abstraction

- Each relation is a file of records.
- Each record has a unique record identifier called RID / TID.
- Common file operations: create/delete file, insert record, delete/get record with given RID, scan all records.

File Organization: Method of arranging data records in a file that is stored on disk.

- **Heap file:** Unordered file
- **Sorted file:** Records order on some search key.
- **Hashed file:** Records located in blocks via a hash function.

Heap File Implementations



- **Linked list implementation:** Two linked lists, one with pages with free space, other of completely full pages.
- **Page Directory Implementation:** Two leveled implementation. Each big block is a disk block with some metadata. Each disk block has a number of data pages.

Page Formats:

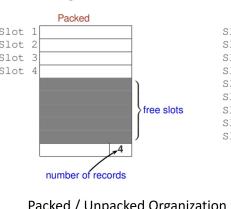
Records are organized within a page and referenced with the RID.

- **RID = (page id, slot number)**
- For **Fixed-Length Records**, Organization can be:

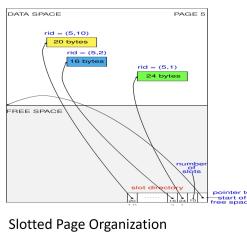
- **Packed Organization:** Store records in contiguous slots.
- For packed organization, memory organization is tough and costly when record in slot is deleted, need to move up a record. But as RID serves as a reference, but need to propagate change in RID.
- **Unpacked Organization:** Uses bit array to maintain free slots.
- For unpacked organization, more bookkeeping needed (use bitmap, 1 & 0 to check if occupied) to store records.

- For **Variable-Length Records**: We could assume some maximum size, then use packed organization. But wasteful. Instead, we can use **Slotted Page Organization**.

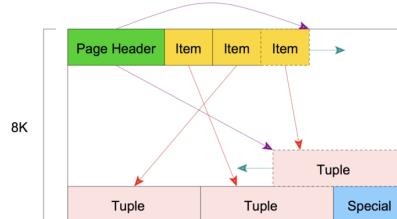
Fixed-Length Records:



Variable-Length Records:



PostgreSQL's Slotted Page Organization



Source: B. Momjian's slides on PostgreSQL internals

Record Formats: Organizing fields within a record.

- **Fixed-Length Records**
 - ▶ Fields are stored consecutively

F1	F2	F3	F4
----	----	----	----
- **Variable-Length Records**
 - ▶ Delimit fields with special symbols
 - ▶ Use an array of field offsets

F1	\$	F2	\$	F3	\$	F4
----	----	----	----	----	----	----

Each o_i is an offset to beginning of field F_i

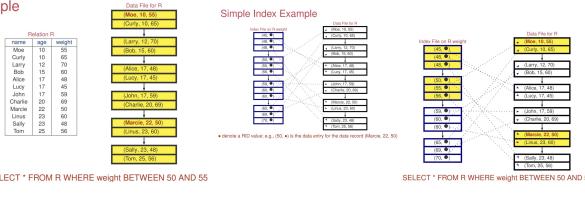
2. Indexing

Need some auxiliary data structure to make efficient queries.

Index

- An **index** is a data structure to speed up retrieval of data records based on some search key.
- A **search key** is a sequence of k data attributes, $k \geq 1$. (A search key is aka *composite search key* if $k > 1$, e.g. (state, city).)
- An index is a **unique index** if search key is a candidate key, otherwise it is **non-unique index**.
- An index is stored as a file, records in index file referred to as **data entries**.

Example



Index Types

Two main types of indexes

- **Tree-based Index:** Based on sorting of search key values (E.g. ISAM, B^+ -tree)
- **Hash-based Index:** Data entries accessed using hashing function (E.g. static/ extendible / linear hashing)
- Considerations when choosing an index:
 - Search Performance (Equality search: $k = v$, use hash-based.) (Range search, use tree)
 - Storage overhead
 - Update performance

Tree-based Indexing: B^+ -Tree

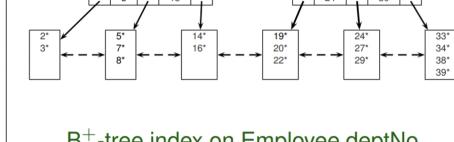
B^+ tree is a dynamic structure that adjusts to changes in the file gracefully, most widely used index structure as it adjusts well to changes and supports both equality and range queries.

- **Balanced tree:** Operations (insert, delete) on tree keep it balanced.
- **Internal nodes** direct the search.
- **Leaf nodes** contain the data entries. Leaf pages linked using page pointers for easy traversal of sequence of leaf pages in either direction.
- **Value d** is parameter of B^+ -tree, called order of the tree, is a measure of capacity of a tree node. Each node contains m entries, where $d \leq m \leq 2d$, except root node, where $1 \leq m \leq 2d$

B^+ -tree Index

B^+ -tree Index

Employee	
name	deptNo
Alice	5
Curly	10
Bob	15
Lucy	17
Cherie	20
Marie	22
Sally	23
Tom	25



- Each node is either a **leaf node** (bottom-most level) or an internal node.
- Top-most internal node is the **root node** located at **level 0**.
- **Height of Tree** = number of level of internal nodes. (Leaf nodes are at level h where $h =$ height of tree).
- Nodes at same level are **sibling nodes** if they have the same parent node.
- **Leaf Nodes:**

- Leaf nodes store sorted data entries.
- $k*$ denote data entry of form (k, RID) , where k = search key value of corresponding data record, $RID =$ RID of data record.
- Lead nodes are doubly-linked to adjacent nodes.
- **Internal Nodes:**
 - Internal nodes store index entries of the form $(p: pointer, k: separator)$ ($p_0, k_1, p_1, k_2, p_2, \dots, p_n$)
 - $k_1 < k_2 < \dots < K_n$
 - Each (k_i, p_i) is an **index entry**, k_i serves as **separator** between node contents pointed to by p_{i-1} & p_i
 - p_i = disk page address (root node of an index subtree T_i)

B^+ -Tree Bulk Loading

- Entries added to a B^+ in two ways.

- Have existing collection of data records with B^+ tree index on it.
When record added to collection, corresponding entry added to B^+ tree. (Insert, Delete individually)
- Have collection of data records we want to create new B^+ tree index on some key field(s). Start with an empty tree. Inserting one by one expensive due to overhead, systems provide **bulk loading utility**.

Bulk Loading:

- Sort data entries $k*$ to be inserted into B^+ tree according to search key k . (Here, $d = 1$).
- Allocate empty page to serve as root. Insert a pointer to first page of (sorted entries into it).

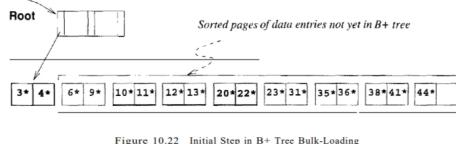


Figure 10.22 Initial Step in B^+ Tree Bulk-Loading

- Add one entry to root page for each page of sorted data entries.
Proceed until root page is full. Here, we must split root and create a new root page.

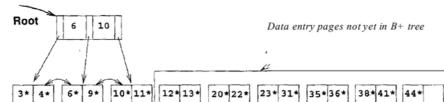


Figure 10.23 Root Page Fills up in B^+ Tree Bulk-Loading

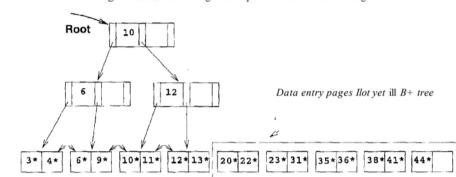


Figure 10.24 Page Split during B^+ Tree Bulk-Loading

- To continue, entries for leaf pages **always inserted into right-most index page just above the leaf level**. When right-most page above leaf level fills up, it is split.

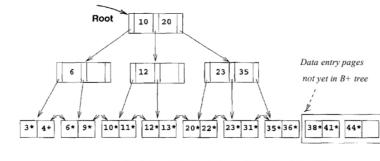


Figure 10.25 Before Adding Entry for Leaf Page Containing 38*

3. Hash-based Indexing

- Used for **equality queries**, not for range queries.
- Hashing techniques:** **Static hashing**, **dynamic hashing** (linear hashing, extendible hashing, etc).

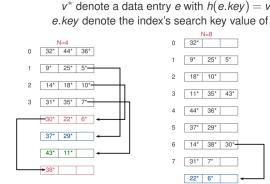
Static Hashing

- Data stored in N buckets**, fixed at creation time. Bucket consists of **one primary data page & chain** of zero or more overflow data pages.
- v^* represents data entry e with $h(e.key) = v$, not search entry with RID.
- Problem with static hashing:** As data grows, longer overflow chain, efficiency drops. Need to periodically rehash and increase no. of buckets.

Static Hashing

- Data is stored in N buckets B_0, B_1, \dots, B_{N-1}
- N is fixed at creation time
- Hashing function $h(\cdot)$ is used to identify the bucket to store a record:
 - $i = h(K) \bmod N$
 - $\{K\}$ maps the search key value into a bit string
- Each bucket consists of **one primary data page** & a chain of zero or more overflow data pages

Static Hashing: Example



Dynamic Hashing: Linear Hashing

- Hash file grows / shrinks linearly, systematic **splitting of buckets**.
- Overflow pages needed as overflowed bucket may not be split immediately.
- Hashing function changes dynamically and at given instant, **at most two (successive) hashing functions** used by the scheme during search.
- Each bucket has primary data page & chain of zero+ overflow pages.
- Insert in bucket B_i overflows if all pages in B_i (primary + overflow) full.

Linear Hashing (cont.)

- Assume initial file size of N_0 buckets
 - Buckets $B_0, B_1, \dots, B_{N_0-1}$
- File grows linearly by **splitting buckets** in rounds
 - At each round, buckets are split sequentially: B_i is split before B_{i+1}
- How to split a bucket B_i ?
 - Add a new bucket B_i' (known as split image of B_i)
 - Redistribute entries in B_i between B_i & B_i'
- File size increases by one bucket after each bucket split
- At the end of one round of splitting (i.e., every bucket at the start of the round has been split), file size is doubled
- Let N_i denote the file size at the beginning of round i ($i=0, 1, \dots$)
 - $N_0 = N_0$
 - At the end of round i , N_i new buckets are added: $B_{N_0}, B_{N_0+1}, \dots, B_{N_i-1}$
 - In round i , the split image of B_i is B_{N_i+i} for $j \in [0, N_i - 1]$

Splitting Buckets

- Assume $N_0 = 4$
 - In round 0, split image of bucket B_0 is B_{0+1} , for $i \in [0, 3]$
- In round 1, $N_1 = 8$ & split image of bucket B_0 is B_{0+4} , for $i \in [0, 7]$

Dynamicity of Linear Hashing

Dynamic Hashing

next = 1



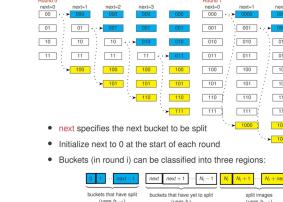
Before B_0 was split in round 0

- Recall that v^* denote a data entry e with $h(e.key) = v$ and e key denote the index's search key value of e
- Each round uses two hash functions: functions h_i and h_{i+1} for round i
 - $h_i(v) = h_{i+1}(v) \bmod N$
 - B_i is the bucket for search key v if B_i had not been split, where $x = h_i(v)$
 - B_i' is the bucket for search key v if B_i had been split, where $y = h_{i+1}(v)$
- Keep track of which bucket to be split next using variable next

Linear (Hashing) Splitting of Buckets (Insertion)

- Number in buckets represent the rightmost bit values.
- Split Criteria:** variable, could be when some bucket overflow, space utilization of file above some threshold etc.
- Level:** We use level to denote splitting round number (use with next).

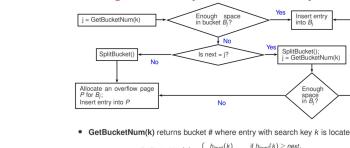
Splitting Buckets



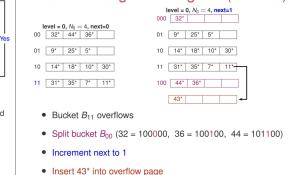
Examples: Linear Hashing Search & Insert

- Even though overflowed bucket split, not necessary mean enough space if bit value not right. Still require overflow page.
- Using **level** and **next**, we determine how many additional bits to consider in **second hashing function**.
- Second hashing function:** Is simply looking at the m rightmost bits of the hash.

Inserting data entry with search key k



Example: Linear Hashing: Inserting 43* (101011)



Dynamic Hashing: Extendible Hashing

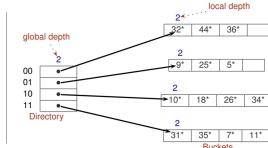
- Similar to Linear Hashing:** we want **bucket number to grow dynamically**, and use some **number of least significant bits** of $h(k)$ to determine bucket address for search key k .
- Difference:** Add new bucket (as split image) when existing bucket overflows. No overflow pages (except when number of collisions exceed page capacity, two page entries collide if they have same $h(.)$ hash value).

Extendible Hashing

- Extendible hashing:** dynamically updateable disk-based index structure, implements hashing scheme utilizing a **directory of pointers to buckets**.
- Overflows handled by doubling the directory which logically doubles the number of buckets. **Physically, only the overflow bucket is split.**

Extendible Hashing

- Uses a directory of pointers to buckets
- Directory expands dynamically as buckets overflow
- Directory has 2^d entries
- d is called global depth of the hashed file
- Each directory entry has a unique d -bit address $b_0 b_{0+1} \dots b_{0+d-1}$
- Two directory entries are said to correspond if their addresses differ only in the i^{th} bit (i.e., b_i). Such entries are called corresponding entries.
- Each bucket maintains a local depth denoted by $\ell \in [0, d]$
- All entries in a bucket with local depth ℓ have the same last ℓ bits in $h(k)$



Extendible Hashing Performance

- Performance:** At most 2 disk I/O for equality selection, at most 1 I/O if directory fits in main memory.
- Handling collision:** Two data entries **collide** if same hashed value, overflow pages need when number of collisions exceed page capacity.
- Compared with B+-tree index exact match queries (log number of I/Os), E. Hashing better expected query cost $O(1)$ I/O.

Extendible Hashing: Handling Bucket Overflow

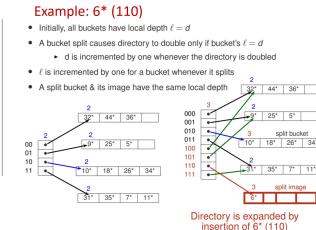
- Main idea: Determine if there is empty directory entry to point to new bucket. **2 Cases:** decision to split, or use empty directory entry.

Case 1: Split bucket local depth = global depth

Extendible Hashing

Handling Bucket Overflow (Case 1)

- When a bucket overflows, it is split
 - Allocate a new bucket called its **split image**
 - Redistribute entries (including new entry) between split bucket & its split image
- Case 1:** Split bucket's local depth is equal to global depth
 - Each new directory entry (except for the entry for the split image) points to the same bucket as its corresponding entry
- Number of directory entries pointing to a bucket = $2^{d-\ell}$

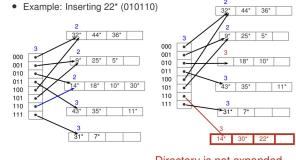


Case 2: Split bucket local depth < global depth.

Extendible Hashing

Handling Bucket Overflow (Case 2)

- Case 2:** Split bucket's local depth < global depth
 - Example: Inserting 22* (010110)



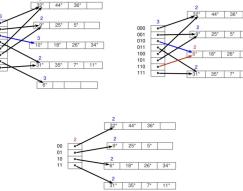
Extendible Hashing Deletion

- To delete entry, simply locate Bucket and delete.
- Merging:** Mergeable if entries can fit within a bucket, and same local depth, j differs on in l^{th} bit.

Extendible Hashing: Deletion

- Locate bucket B_j containing entry & delete entry
- If B_j becomes empty, B_j can be merged with the bucket B_i where both buckets have the same local depth ℓ and $i \neq j$ differs only in the l^{th} bit
 - B_i is deallocated
 - B_j 's local depth is decremented by one
 - Directory entries that point to B_i are updated to point to B_j
- More generally, B_i & B_j (with same local depth ℓ and $i \neq j$ differs only in the l^{th} bit) can be merged if their entries can fit within a bucket
- If each pair of corresponding entries point to the same bucket, directory can be halved
 - d is decremented by one

Example: Deleting 10* (1010)



Extendible Hashing

Both cases: Inserting data entry with search key k

