

Files

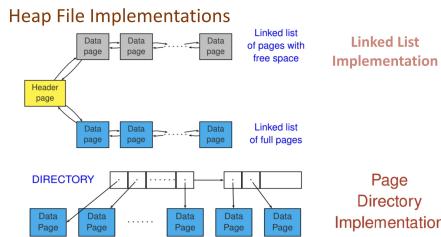
File Abstraction

- Each relation is a file of records.
- Each record has a unique record identifier called RID / TID.
- Common file operations: create/delete file, insert record, delete/get record with given RID, scan all records.

File Organization: Method of arranging data records in a file that is stored on disk.

- **Heap file:** Unordered file
- **Sorted file:** Records order on some search key.
- **Hashed file:** Records located in blocks via a hash function.

Heap File Implementations



- **Linked list implementation:** Two linked lists, one with pages with free space, other of completely full pages.
- **Page Directory Implementation:** Two leveled implementation. Each big block is a disk block with some metadata. Each disk block has a number of data pages.

Page Formats:

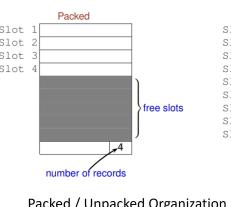
Records are organized within a page and referenced with the RID.

- **RID = (page id, slot number)**
- For **Fixed-Length Records**, Organization can be:

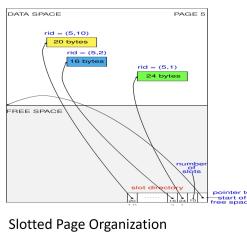
- **Packed Organization:** Store records in contiguous slots.
- For packed organization, memory organization is tough and costly when record in slot is deleted, need to move up a record. But as RID serves as a reference, but need to propagate change in RID.
- **Unpacked Organization:** Uses bit array to maintain free slots.
- For unpacked organization, more bookkeeping needed (use bitmap, 1 & 0 to check if occupied) to store records.

- For **Variable-Length Records**: We could assume some maximum size, then use packed organization. But wasteful. Instead, we can use **Slotted Page Organization**.

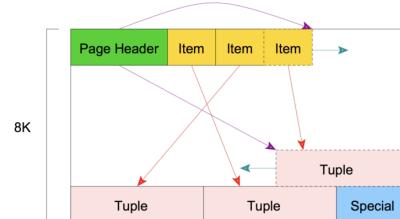
Fixed-Length Records:



Variable-Length Records:



PostgreSQL's Slotted Page Organization



Source: B. Momjian's slides on PostgreSQL internals

Record Formats: Organizing fields within a record.

- **Fixed-Length Records**
 - ▶ Fields are stored consecutively

F1	F2	F3	F4
----	----	----	----
- **Variable-Length Records**
 - ▶ Delimit fields with special symbols
 - ▶ Use an array of field offsets

F1	\$	F2	\$	F3	\$	F4
----	----	----	----	----	----	----

Each o_i is an offset to beginning of field F_i

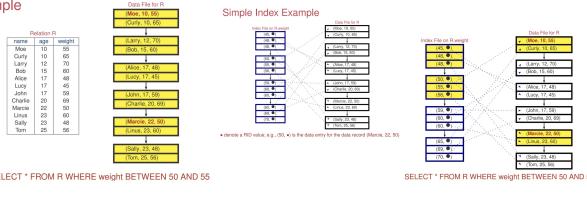
2. Indexing

Need some auxiliary data structure to make efficient queries.

Index

- An **index** is a data structure to speed up retrieval of data records based on some search key.
- A **search key** is a sequence of k data attributes, $k \geq 1$. (A search key is aka *composite search key* if $k > 1$, e.g. (state, city).)
- An index is a **unique index** if search key is a candidate key, otherwise it is **non-unique index**.
- An index is stored as a file, records in index file referred to as **data entries**.

Example



w/o index, need to retrieve all pages/records

w. simple index, only require retrieval of select pages

Index Types

Two main types of indexes

- **Tree-based Index:** Based on sorting of search key values (E.g. ISAM, B^+ -tree)
- **Hash-based Index:** Data entries accessed using hashing function (E.g. static/ extendible / linear hashing)
- Considerations when choosing an index:
 - Search Performance (Equality search: $k = v$, use hash-based.) (Range search, use tree)
 - Storage overhead
 - Update performance

B^+ -Tree

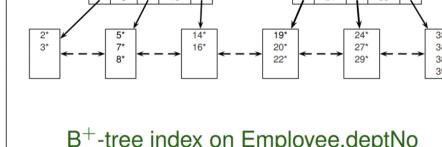
B^+ tree is a dynamic structure that adjusts to changes in the file gracefully, most widely used index structure as it adjusts well to changes and supports both equality and range queries.

- **Balanced tree:** Operations (insert, delete) on tree keep it balanced.
- **Internal nodes** direct the search.
- **Leaf nodes** contain the data entries. Leaf pages linked using page pointers for easy traversal of sequence of leaf pages in either direction.
- **Value d** is parameter of B^+ -tree, called order of the tree, is a measure of capacity of a tree node. Each node contains m entries, where $d \leq m \leq 2d$, except root node, where $1 \leq m \leq 2d$

B^+ -tree Index

B^+ -tree Index

Employee	
name	deptNo
Alice	5
Curly	19
Bob	39
Dave	38
Eve	14
Fred	33
Harry	2
John	34
Ken	6
Larry	27
Linus	24
Lucy	3
Marcie	22
Moe	29
Sally	20
Tom	7
	...



- Each node is either a **leaf node** (bottom-most level) or an internal node.
- Top-most internal node is the **root node** located at **level 0**.
- **Height of Tree** = number of level of internal nodes. (Leaf nodes are at level h where $h =$ height of tree).
- Nodes at same level are **sibling nodes** if they have the same parent node.
- **Leaf Nodes:**

- Leaf nodes store sorted data entries.
- $k*$ denote data entry of form (k, RID) , where k = search key value of corresponding data record, $RID =$ RID of data record.
- Lead nodes are doubly-linked to adjacent nodes.

- **Internal Nodes:**

- Internal nodes store index entries of the form $(p: pointer, k: separator)$ ($p_0, k_1, p_1, k_2, p_2, \dots, p_n$)
- $k_1 < k_2 < \dots < K_n$
- Each (k_i, p_i) is an **index entry**, k_i serves as **separator** between node contents pointed to by p_{i-1} & p_i
- p_i = disk page address (root node of an index subtree T_i)

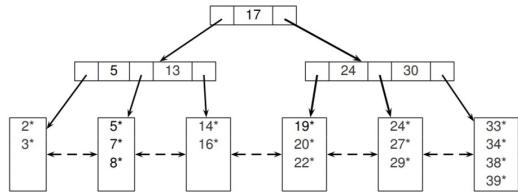
B^+ -tree Index Properties

Properties of B^+ -tree Index

- Dynamic index structure; adapts to data updates gracefully
- Height-balanced index structure
- Order of index tree, $d \in \mathbb{Z}^+$

1. Controls space utilization of index nodes
2. Each non-root node contains m entries, where $m \in [d, 2d]$
3. The root node contains m entries, where $m \in [1, 2d]$

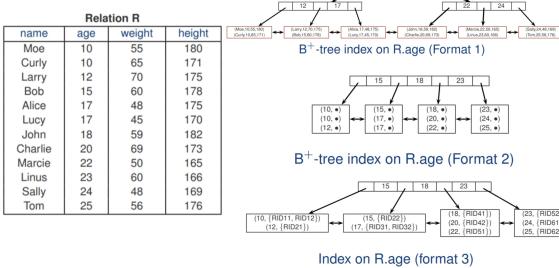
Example: B^+ -tree with order = 2



Formats of Data Entries in B-Tree

- Format 1: k^* is actual data record (with search key value k)
- Format 2: k^* is of form (k, rid) , where rid is record identifier of record with search key value k .
- Format 3: k^* is of form $(k, rid-list)$, where rid-list is list of record identifiers of data records with search key value k .
- Note, examples assume Format 2.

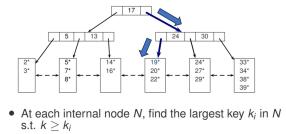
Formats of Data Entries: Example



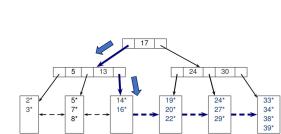
B^+ -tree Search Algorithms

- Search algorithm finds the leaf node a given data entry belongs to.
- We assume no duplicates, no data entries same key value. Note in practice, duplicates arise whenever search key does not contain candidate key, must be dealt with.

Equality Search ($k = 19$)



Range Search ($k \geq 15$)



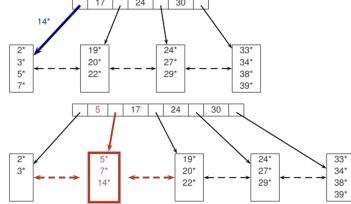
B^+ -Tree Insertion

- Algorithm for insertion takes an entry, finds the leaf node where it belongs, and inserts it there.
- Occasionally, a node is full and must be split. (More than $2d$ entries)
 When node is split, entry pointing to the node created by the split must be inserted into the parent.
- If the (old) root is split, a new root node is created and height of tree increases by 1.

Splitting of overflowed node

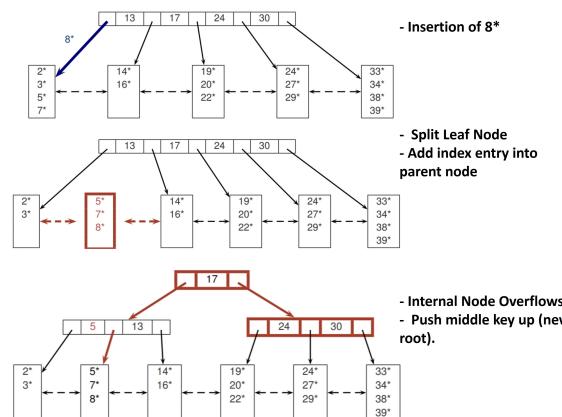
- Split overflowed leaf node by distributing $d + 1$ entries to new leaf node.
- Create a new entry index using smallest key in leaf node.
- Insert new index entry into parent node of overflowed node.

Inserting 14* (Splitting of overflowed node)

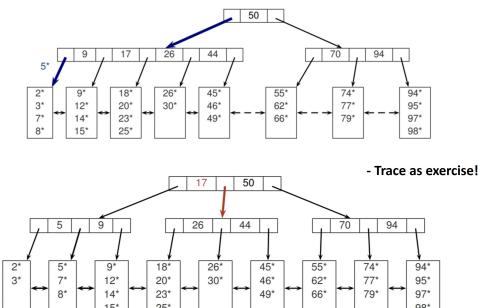


- Sometimes, node split is propagated upwards to ancestor internal nodes.
- When splitting an internal node, the middle key is pushed to parent node.

Inserting 8* (Propagation of node splits)



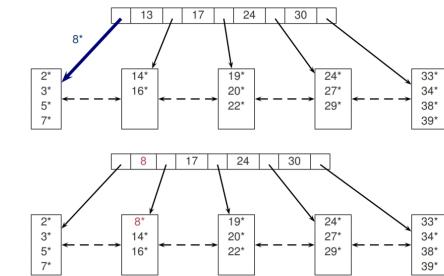
Inserting 5* (Propagation of node splits)



Redistributing of data entries in Overflow

- A node split can sometimes be avoided by distributing entries from overflowed node to a non-full adjacent sibling node.

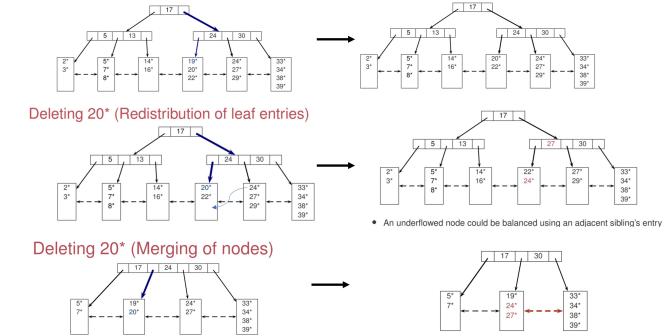
Inserting 8* (Redistribution of data entries)



B^+ -Tree Deletion

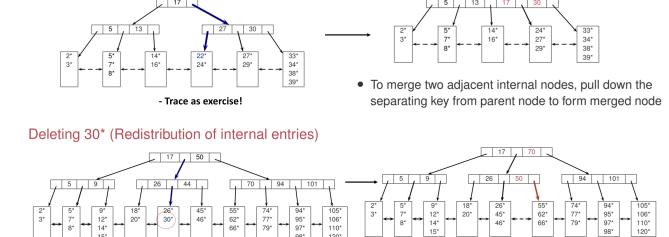
- Algorithm for deletion takes an entry, finds leaf node it belongs to, and deletes it.
- Underflowed node: When node is at minimum occupancy before deletion, and goes below threshold, we must either redistribute entries from adjacent sibling, or merge node with sibling to maintain minimum occupancy.
- Merging: Underflowed node needs to be merged if each of adjacent sibling nodes has exactly d entries.

Deleting 19* (Simple Case)

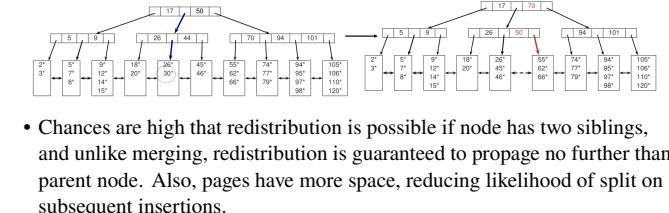


- Node mergers may propagate upwards.

Deleting 22* (Propagation of node merges)



Deleting 30* (Redistribution of internal entries)



B^+ -Tree Bulk Loading

- Entries added to a B^+ in two ways.

- Have existing collection of data records with B^+ tree index on it.
When record added to collection, corresponding entry added to B^+ tree. (Insert, Delete individually)
- Have collection of data records we want to create new B^+ tree index on some key field(s). Start with an empty tree. Inserting one by one expensive due to overhead, systems provide **bulk loading** utility.

Bulk Loading:

- Sort data entries $k*$ to be inserted into B^+ tree according to search key k .
- Allocate empty page to serve as root. Insert a pointer to first page of (sorted entries into it).



Figure 10.22 Initial Step in B+ Tree Bulk-Loading

- Add one entry to root page for each page of sorted data entries.
Proceed until root page is full. Here, we must split root and create a new root page.

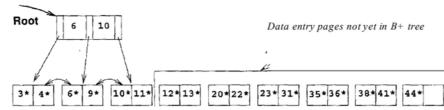


Figure 10.23 Root Page Fills up in B+ Tree Bulk-Loading

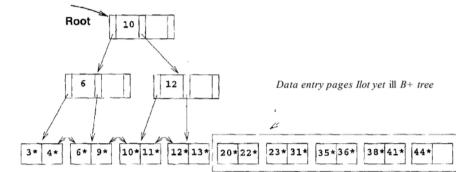


Figure 10.24 Page Split during B+ Tree Bulk-Loading

- To continue, entries for leaf pages **always inserted into right-most index page just above the leaf level**. When right-most page above leaf level fills up, it is split.

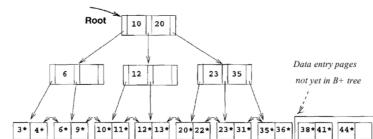


Figure 10.25 Before Adding Entry for Leaf Page Containing 38*

