

CS3223 Database Systems

Implementation

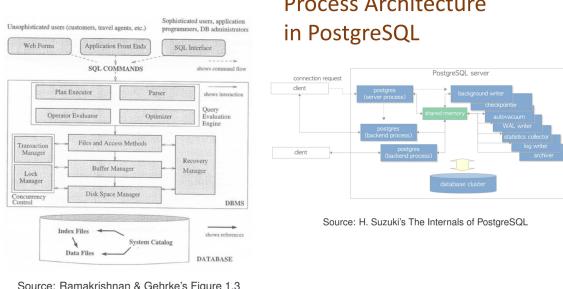
AY23/24 Sem 2, github.com/gerteck

Introduction

Course Details

- Prerequisite knowledges: CS2040S, CS2102, CS2106 background (helpful).
- Reference Textbook: Raghu & Johannes Database M. Systems, 2002. Encouraged to read ahead based on schedule before the lecture.
- Course covers data structures, algorithms, different components making up database systems.

Architecture of DBMS



Source: Ramakrishnan & Gehrke's Figure 1.3

- OLTP:** Online Transaction Processing is a type of data processing that consists of executing a number of transactions occurring concurrently—online banking, shopping, order entry, or sending text messages, for example.
- OLAP:** Online Analytical Processing.
- Focusing on centralized database running on a single server.

1. Data Storage

References: R&G Chapt 8. (Storage & Indexing Overview), Chapt 9. (Storing Data: Disks and Files).

A DBMS stores

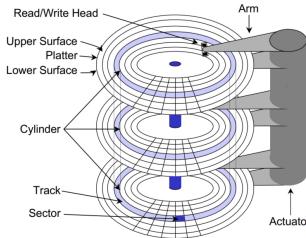
- Relations (Actual tables)
- System catalog (aka data dictionary) storing metadata about relations. (Relation schemas, structure of relations, constraints, triggers. View definitions, Indexes - derived info to speed up access to relations, Statistical information about relations for use by query optimizer.)
- Log files: Information maintained for data recovery.

DBMS Storage

Memory Hierarchy: Primary (registers, RAM), secondary (HDD, SSD), tertiary memory with capacity / cost / access speed / volatility tradeoffs.

- DBMS stores data on non-volatile disk for persistence.
- DBMS processes data in main memory (RAM).
- Disk access operations (I/O). Read: transfer data from disk to RAM. Write: transfer data from RAM to disk.
- Make use of index to speed up access, so that don't have to retrieve all the data when you run a query. Retrieve index and read only the block that contains specified data. Minimize I/O cost.

Magnetic Hard-Disk Drive HDD



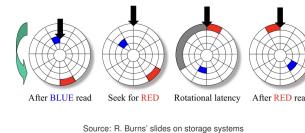
Source: R. Burns' slides on storage systems

- Cylinder, Track, Sector: Units of the HDD storage system. To read from different tracks, need to move the mechanical HDD arm.

Disk Access Time:

- command processing time: interpreting access command by disk controller.
- seek time: moving arms to position disk head on track.
- rotational delay: waiting for block to rotate under head.
- transfer time: actually moving data to/from disk surface.
- access time:** = seek time + rotational delay + transfer time. (CPT considered negligible).

Disk Access Time Components



- Seek time
 - avg. seek time: 5-6 ms
- Rotational delay (or rotational latency)
 - Depends on rotation speed - measured in rotations per minute (RPM)
 - Average rotational delay = time for $\frac{1}{2}$ revolution
 - Example: For 10000 RPM, avg. rotational delay = $0.5 \times (60 / 10000) = 3$ ms
- Transfer time
 - n = number of requested sectors on same track
 - transfer time = $n \times \text{time for one revolution}$
 - avg. sector transfer time: 100-200 μ s
 - Sequential vs random I/O

Concept of Sequential vs random I/O.

- Sequential:** Both sector on same track.
- Random:** Sectors on different track, require seeking (moving arm).
- Given a set of data, we hope to store the data contiguously, on the same track. (Minimize incurring random I/O). If data is too large, store on same track, but different surface (aka same cylinder).
- Complexity hidden to OS by disk controller. Shown as a sequence of memory locations.

Solid-State Drive: SSD

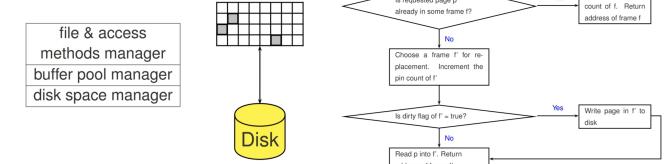
- Build with NAND flash memory without any mechanical moving parts. Lower power consumption.
- Random I/O:** 100x faster than HDD. (no moving parts)
- Sequential I/O:** slightly faster than HDD (2x)
- Disadvantages:** update to a page requires erasure of multiple pages (5ms) before overwriting page. Limited number of times a page can be erased ($10^5 - 10^6$)

Storage Manager Components

- Data is stored, retrieved in units called **disk blocks (or pages)**. (Each block = sequence of one or more contiguous sectors).
- Files & access methods layer (aka file layer)** - deals with organization and retrieval of data.
- Buffer Manager** - controls reading/writing of disk pages.

- Disk Space Manager** - keeps track of pages used by file layer.

Storage Manager Buffer Manager BM: Handling request for page p Components



Buffer Manager

- Buffer pool:** Main memory allocated for DBMS.
- Buffer pool is partitioned into block-sized pages called **frames**.
- Clients of buffer pool can request for disk page to be fetched into buffer pool.
- A page in the buffer is **dirty** if it has been modified & not updated on disk.
- Two variables** maintained for each frame in buffer pool:
 - pin count:** number of clients using page (initialized 0)
 - dirty flag:** whether page is dirty (initialized false)
- Free list: Keeps track of frames that are free / empty.
- Pin count:**
 - Incrementing pin count is **pinning** the requested page in its frame.
 - Decrementing is **unpinning** the page.
- Unpinning a page, dirty flag should be updated to true if page is dirty.
- A page in buffer can be replaced only when pin count is 0.
- Before replacing buffer page, needs to be written back to disk if its dirty flag is true.

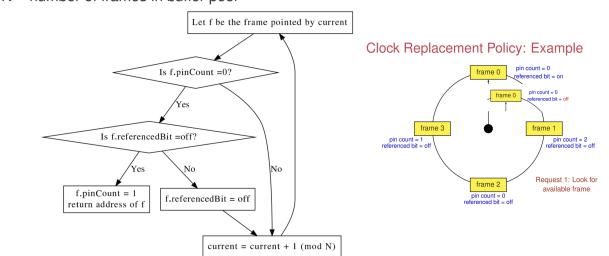
- Buffer manager coordinates with transaction manager to ensure data correctness and recoverability.

Replacement Policies

- Replacement policy: Deciding which unpinned page to replace. (some examples:)
- Random, FIFO, Most Recently Used (MRU), Least Recently Used (LRU): (Use queue of pointers to frames with pin count = 0), most common, makes use of temporal locality.
- Clock:** cheaper popular variant of LRU
 - current** variable: points to some buffer frame.
 - Each frame has a **referenced bit**, turns on when its pin count turns 0.
 - Replace a page that has referenced bit off & pin count = 0.

Clock Replacement Policy

N = number of frames in buffer pool



CS3223: Sem 2, 2023/24

Files

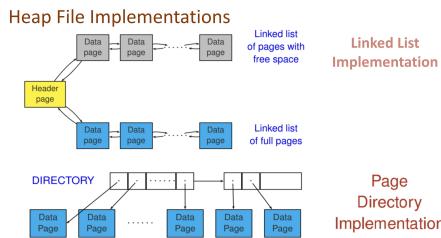
File Abstraction

- Each relation is a file of records.
- Each record has a unique record identifier called RID / TID.
- Common file operations: create/delete file, insert record, delete/get record with given RID, scan all records.

File Organization: Method of arranging data records in a file that is stored on disk.

- **Heap file:** Unordered file
- **Sorted file:** Records order on some search key.
- **Hashed file:** Records located in blocks via a hash function.

Heap File Implementations



- **Linked list implementation:** Two linked lists, one with pages with free space, other of completely full pages.
- **Page Directory Implementation:** Two leveled implementation. Each big block is a disk block with some metadata. Each disk block has a number of data pages.

Page Formats:

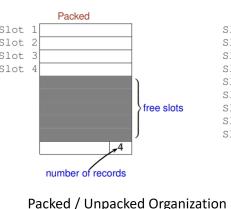
Records are organized within a page and referenced with the RID.

- **RID = (page id, slot number)**
- For **Fixed-Length Records**, Organization can be:

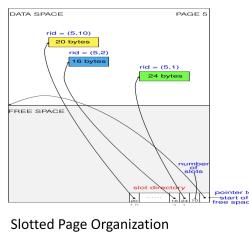
- **Packed Organization:** Store records in contiguous slots.
- For packed organization, memory organization is tough and costly when record in slot is deleted, need to move up a record. But as RID serves as a reference, but need to propagate change in RID.
- **Unpacked Organization:** Uses bit array to maintain free slots.
- For unpacked organization, more bookkeeping needed (use bitmap, 1 & 0 to check if occupied) to store records.

- For **Variable-Length Records**: We could assume some maximum size, then use packed organization. But wasteful. Instead, we can use **Slotted Page Organization**.

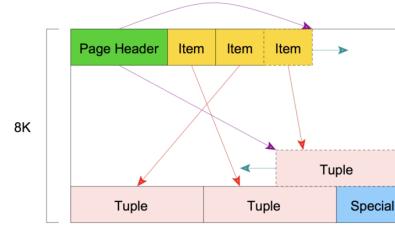
Fixed-Length Records:



Variable-Length Records:



PostgreSQL's Slotted Page Organization



Source: B. Momjian's slides on PostgreSQL internals

Record Formats: Organizing fields within a record.

- **Fixed-Length Records**
 - ▶ Fields are stored consecutively

F1	F2	F3	F4
----	----	----	----
- **Variable-Length Records**
 - ▶ Delimit fields with special symbols
 - ▶ Use an array of field offsets

F1	\$	F2	\$	F3	\$	F4
----	----	----	----	----	----	----

Each o_i is an offset to beginning of field F_i

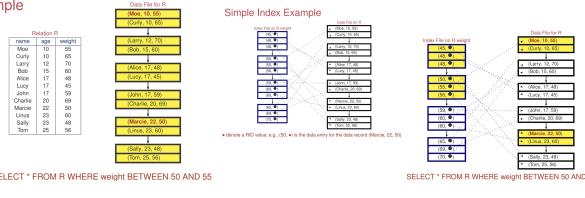
2. Indexing

Need some auxiliary data structure to make efficient queries.

Index

- An **index** is a data structure to speed up retrieval of data records based on some search key.
- A **search key** is a sequence of k data attributes, $k \geq 1$. (A search key is aka *composite search key* if $k > 1$, e.g. (state, city).)
- An index is a **unique index** if search key is a candidate key, otherwise it is **non-unique index**.
- An index is stored as a file, records in index file referred to as **data entries**.

Example



w/o index, need to retrieve all pages/records

w. simple index, only require retrieval of select pages

Index Types

Two main types of indexes

- **Tree-based Index:** Based on sorting of search key values (E.g. ISAM, B^+ -tree)
- **Hash-based Index:** Data entries accessed using hashing function (E.g. static/ extendible / linear hashing)
- Considerations when choosing an index:
 - Search Performance (Equality search: $k = v$, use hash-based.) (Range search, use tree)
 - Storage overhead
 - Update performance

Tree-based Indexing: B^+ -Tree

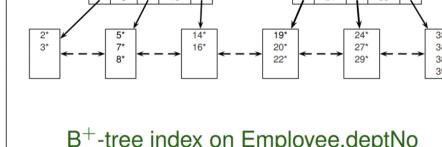
B^+ tree is a dynamic structure that adjusts to changes in the file gracefully, most widely used index structure as it adjusts well to changes and supports both equality and range queries.

- **Balanced tree:** Operations (insert, delete) on tree keep it balanced.
- **Internal nodes** direct the search.
- **Leaf nodes** contain the data entries. Leaf pages linked using page pointers for easy traversal of sequence of leaf pages in either direction.
- **Value d** is parameter of B^+ -tree, called order of the tree, is a measure of capacity of a tree node. Each node contains m entries, where $d \leq m \leq 2d$, except root node, where $1 \leq m \leq 2d$

B^+ -tree Index

B^+ -tree Index

Employee	
name	deptNo
Alice	5
Bob	17
Charlie	19
Curly	39
Dave	38
Eve	14
Fred	33
Harry	2
John	34
Ken	6
Larry	27
Linus	24
Lucy	3
Marcie	22
Moe	29
Sally	20
Tom	7



- Each node is either a **leaf node** (bottom-most level) or an internal node.
- Top-most internal node is the **root node** located at **level 0**.
- **Height of Tree** = number of level of internal nodes. (Leaf nodes are at level h where $h =$ height of tree).
- Nodes at same level are **sibling nodes** if they have the same parent node.
- **Leaf Nodes:**

- Leaf nodes store sorted data entries.
- k^* denote data entry of form (k, RID) , where k = search key value of corresponding data record, $RID =$ RID of data record.
- Lead nodes are doubly-linked to adjacent nodes.

• Internal Nodes:

- Internal nodes store index entries of the form $(p: pointer, k: separator)$ ($p_0, k_1, p_1, k_2, p_2, \dots, p_n$)
- $k_1 < k_2 < \dots < K_n$
- Each (k_i, p_i) is an **index entry**, k_i serves as **separator** between node contents pointed to by p_{i-1} & p_i
- p_i = disk page address (root node of an index subtree T_i)

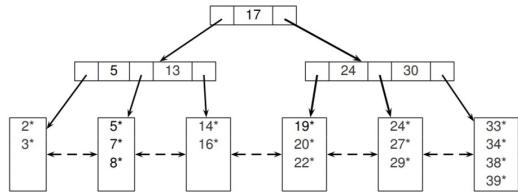
B^+ -tree Index Properties

Properties of B^+ -tree Index

- Dynamic index structure; adapts to data updates gracefully
- Height-balanced index structure
- Order of index tree, $d \in \mathbb{Z}^+$

1. Controls space utilization of index nodes
2. Each non-root node contains m entries, where $m \in [d, 2d]$
3. The root node contains m entries, where $m \in [1, 2d]$

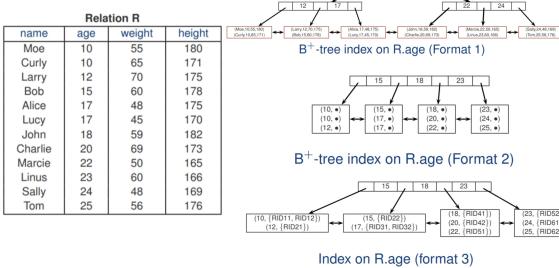
Example: B^+ -tree with order = 2



Formats of Data Entries in B-Tree

- Format 1: k^* is actual data record (with search key value k)
- Format 2: k^* is of form (k, rid) , where rid is record identifier of record with search key value k .
- Format 3: k^* is of form $(k, rid-list)$, where rid-list is list of record identifiers of data records with search key value k .
- Note, examples assume Format 2.

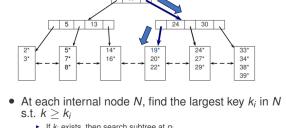
Formats of Data Entries: Example



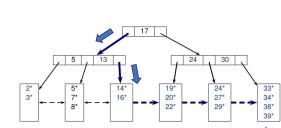
B^+ -tree Search Algorithms

- Search algorithm finds the leaf node a given data entry belongs to.
- We assume no duplicates, no data entries same key value. Note in practice, duplicates arise whenever search key does not contain candidate key, must be dealt with.

Equality Search ($k = 19$)



Range Search ($k \geq 15$)



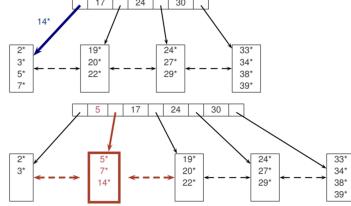
B^+ -Tree Insertion

- Algorithm for insertion takes an entry, finds the leaf node where it belongs, and inserts it there.
- Occasionally, a node is full and must be split. (More than $2d$ entries)
 When node is split, entry pointing to the node created by the split must be inserted into the parent.
- If the (old) root is split, a new root node is created and height of tree increases by 1.

Splitting of overflowed node

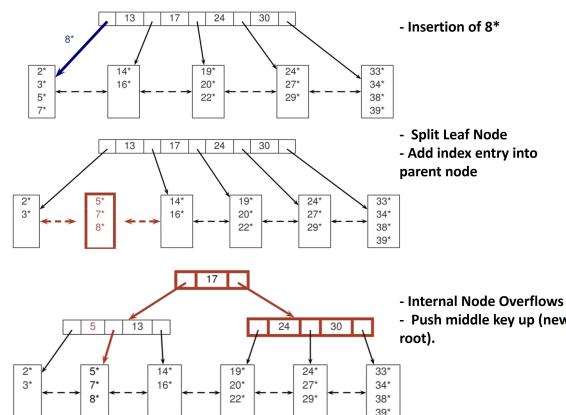
- Split overflowed leaf node by distributing $d + 1$ entries to new leaf node.
- Create a new entry index using smallest key in leaf node.
- Insert new index entry into parent node of overflowed node.

Inserting 14* (Splitting of overflowed node)

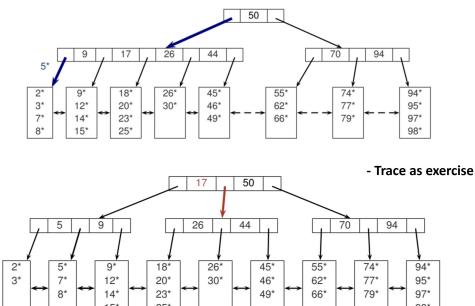


- Sometimes, node split is propagated upwards to ancestor internal nodes.
- When splitting an internal node, the middle key is pushed to parent node.

Inserting 8* (Propagation of node splits)



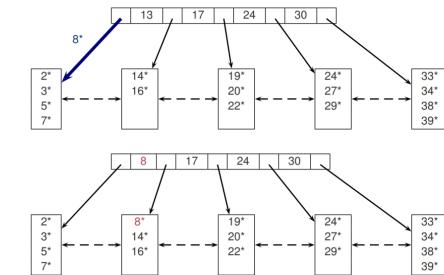
Inserting 5* (Propagation of node splits)



Redistributing of data entries in Overflow

- A node split can sometimes be avoided by distributing entries from overflowed node to a non-full adjacent sibling node.

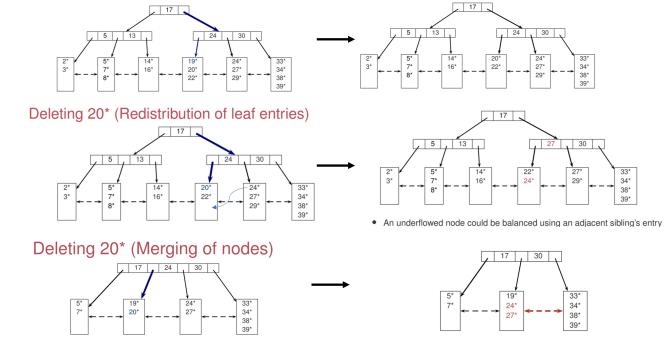
Inserting 8* (Redistribution of data entries)



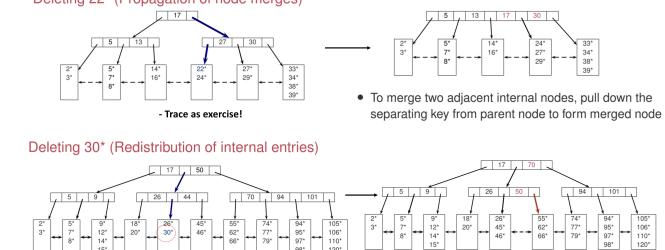
B^+ -Tree Deletion

- Algorithm for deletion takes an entry, finds leaf node it belongs to, and deletes it.
- Underflowed node: When node is at minimum occupancy before deletion, and goes below threshold, we must either redistribute entries from adjacent sibling, or merge node with sibling to maintain minimum occupancy.
- Merging: Underflowed node needs to be merged if each of adjacent sibling nodes has exactly d entries.

Deleting 19* (Simple Case)



Deleting 22* (Propagation of node merges)



Deleting 30* (Redistribution of internal entries)

- Chances are high that redistribution is possible if node has two siblings, and unlike merging, redistribution is guaranteed to propagate no further than parent node. Also, pages have more space, reducing likelihood of split on subsequent insertions.

B⁺ Tree Bulk Loading

- Entries added to a B⁺ in two ways.

- Have existing collection of data records with B⁺ tree index on it.
When record added to collection, corresponding entry added to B⁺ tree. (Insert, Delete individually)
- Have collection of data records we want to create new B⁺ tree index on some key field(s). Start with an empty tree. Inserting one by one expensive due to overhead, systems provide **bulk loading utility**.

Bulk Loading:

- Sort data entries $k*$ to be inserted into B⁺ tree according to search key k . (Here, $d = 1$).
- Allocate empty page to serve as root. Insert a pointer to first page of (sorted entries into it).

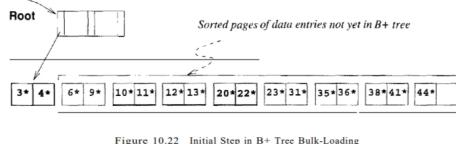


Figure 10.22 Initial Step in B+ Tree Bulk-Loading

- Add one entry to root page for each page of sorted data entries.
Proceed until root page is full. Here, we must split root and create a new root page.

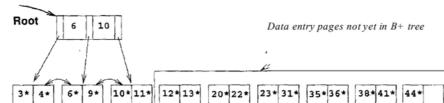


Figure 10.23 Root Page Fills up in B+ Tree Bulk-Loading

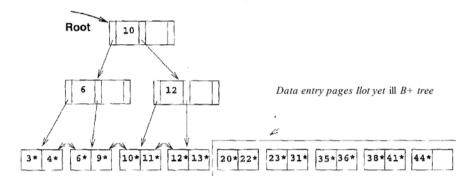


Figure 10.24 Page Split during B+ Tree Bulk-Loading

- To continue, entries for leaf pages **always inserted into right-most index page just above the leaf level**. When right-most page above leaf level fills up, it is split.

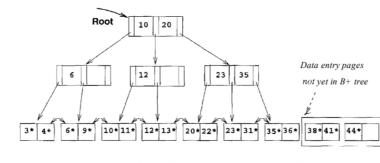
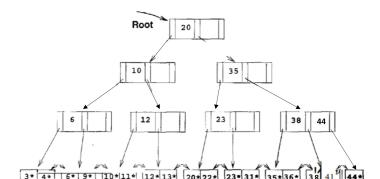


Figure 10.25 Before Adding Entry for Leaf Page Containing 38*



3. Hash-based Indexing

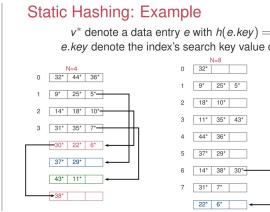
- Used for **equality queries**, not for range queries.
- Hashing techniques:** **Static hashing, dynamic hashing** (linear hashing, extendible hashing, etc).

Static Hashing

- Data stored in N buckets**, fixed at creation time. Bucket consists of **one primary data page & chain** of zero or more overflow data pages.
- v^* represents data entry e with $h(e.key) = v$, not search entry with RID.
- Problem with static hashing:** As data grows, longer overflow chain, efficiency drops. Need to periodically rehash and increase no. of buckets.

Static Hashing

- Data is stored in N buckets B_0, B_1, \dots, B_{N-1}
- N is fixed at creation time
- Hashing function $h(\cdot)$ is used to identify the bucket to store a record
 - A record with search key K is inserted into bucket B_i , where $i = K \mod N$
 - (K, A) maps the search key value to a bit string
- Each bucket consists of **one primary data page** & a chain of zero or more overflow data pages



Dynamic Hashing: Linear Hashing

- Hash file grows / shrinks linearly, systematic **splitting of buckets**.
- Overflow pages needed as overflowed bucket may not be split immediately.
- Hashing function changes dynamically and at given instant, **at most two (successive) hashing functions** used by the scheme during search.
- Each bucket has primary data page & chain of zero+ overflow pages.
- Insert in bucket B_i overflows if all pages in B_i (primary + overflow) full.

Linear Hashing (cont.)

- Assume initial file size of N_0 buckets
 - Buckets $B_0, B_1, \dots, B_{N_0-1}$
- File grows linearly by **splitting buckets in rounds**
 - At each round, buckets are split sequentially: B_i is split before B_{i+1}
- How to split a bucket B_i ?
 - Add a new bucket B_i' (known as split image of B_i)
 - Redistribute entries in B_i between B_i & B_i'
- File size increases by one bucket after each bucket split
- At the end of one round of splitting (i.e., every bucket at the start of the round has been split), file size is doubled.
- Let N_i denote the file size at the beginning of round i ($i = 0, 1, \dots$)
 - $N_0 = N_0$
 - At the end of round i , N_i new buckets are added: $B_{N_0}, B_{N_0+1}, \dots, B_{N_i-1}$
 - In round i , the split image of B_i is B_{N_i+i} for $j \in [0, N_i - 1]$
- Before B_i was split in round 0

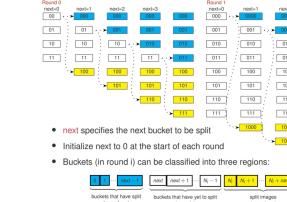
After B_i was split in round 0

- For convenience, $N_i = 2^m$ for $m \in \{0, 1, \dots\}$
- $N_i = 2^{N_0} = 2^{2^{i-1}}$
- B_i is the bucket for search key K if B_i had not been split, where $x = h(K)$
- B_i is the bucket for search key K if B_i had been split, where $y = h_i(x)$
- Keep track of which bucket to be split next using variable $next$

Linear (Hashing) Splitting of Buckets (Insertion)

- Number in buckets represent the rightmost bit values.
- Split Criteria:** variable, could be when some bucket overflow, space utilization of file above some threshold etc.
- Level:** We use level to denote splitting round number (use with $next$).

Splitting Buckets



Examples: Linear Hashing Search & Insert

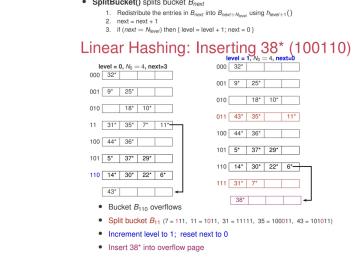
- Even though overflowed bucket split, not necessary mean enough space if bit value not right. Still require overflow page.
- Using **level** and **next**, we determine how many additional bits to consider in **second hashing function**.
- Second hashing function:** Is simply looking at the m rightmost bits of the hash.

Inserting data entry with search key k



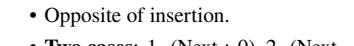
- Bucket B_{N_0} overflows
- Split bucket B_{N_0} ($N_0 = 1000000$, $32 = 1000000$, $36 = 1001000$, $44 = 1011000$)
- Increment next to 1
- Insert 43^* into overflow page

Linear Hashing: Inserting 38* (100110)



- Bucket B_{N_0} overflows
- Split bucket B_{N_0} ($N_0 = 1000000$, $32 = 1000000$, $36 = 1001000$, $44 = 1011000$)
- Increment level to 1; reset next to 0
- Insert 38^* into overflow page

Linear Hashing: Inserting 62* (110110)



- No overflow, no splitting
- Insert 62^* into bucket B_{N_0}

Linear Hashing Deletion

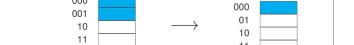
- Opposite of insertion.

- Two cases:** 1. (Next $\neq 0$), 2. (Next = 0, and level > 0.)

Linear Hashing: Deletion

- Locate bucket and delete entry
- If the last bucket $B_{N_{level-1}+next-1}$ becomes empty, it can be removed
- Case 1:** If $next > 0$
 - Decrement next by one

level = 1, $N_0 = 4$, $next=2$



- level = 1, $N_0 = 2$, $next = 2$



- level = 1, $N_0 = 2$, $next = 1$

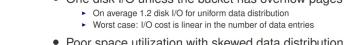


- level = 1, $N_0 = 1$, $next = 1$

Linear Hashing: Deletion (cont.)

- Case 2:** If $(next = 0)$ and $(level > 0)$
 - Update $next$ to point to the last bucket in previous level $B_{N_{level-1}-1}$
 - Decrement $level$ by one

level = 1, $N_0 = 1$, $next=0$



- level = 0, $N_0 = 1$, $next = 1$

Linear Hashing: Summary

- File grows dynamically by splitting buckets linearly in rounds
 - Each bucket split adds one new bucket to file
- The bucket B_i for a search key K is determined using two hash functions: h_i and h_{i+1} for round i
 - $h_i(K) = h(K) \bmod N_i$
 - B_i is the bucket for search key K if B_i had not been split, where $x = h_i(K)$
 - B_i is the bucket for search key K if B_i had been split, where $y = h_{i+1}(K)$
 - Consider the splitting of bucket B_i in round i
 - Before splitting, all entries in B_i have $h_i(e.key) = j$ (i.e., same last $(m-i)$ bits of $h(e.key)$)
 - After splitting, e remains in B_i if the last $(m-i)$ bit of $h(e.key)$ is 0

• Assume initial file has 2^n buckets & $N_0 = 2^n$. If N_0 is not yet been split in current round, then $i = 1$ (i.e., $m = N_0$ bits of $H(K)$)

• In this course, we assume a bucket split is triggered whenever some bucket overflows due to an insertion

• Overflow pages are required for an overflown bucket if it is not the next bucket to be split

External Sort Optimization: I/O Cost vs. No. of I/Os

- Cost Metric:** No. of page I/Os.
- Only an approx of true I/O costs, ignores effect of **block(ed)** I/O, **single request to read/write several consecutive pages can be cheaper** than read/write same number of pages through independent I/O requests.
- Others:** Consider CPU costs as well, can use *double buffering* to keep CPU busy while I/O op. in progress.

External Sort, Block(ed) Page I/O Optimization

Non-Block(ed) Page I/O

- Consider only No. of page I/O as metric:** Minimize no. of passes in sorting, as each page in file read and written in each pass.
- This means we maximise fan-in **during merging** (aka, how many runs merged per pass), allocate just one buffer pool page per run, and one buffer page for output of merge.
- Hence, (B-1)-way merge per run, minimizing number of passes in sorting algorithm.

Block(ed) Page I/O

- Optimization:** Access in blocks may reduce average cost to read/write single page, consideration to read/write in units of more than one page, or R/W in units of a **buffer block**.
- Buffer block**, of b pages. B is total number of buffer pages for sorting.

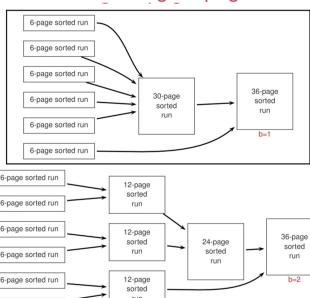
Optimization with Blocked I/O

- Read and write in units of **buffer blocks** of b pages
- Given an allocation of B buffer pages for sorting,
 - Allocate one block (b pages) for output
 - Remaining space can accommodate $\lceil \frac{B-b}{b} \rceil$ blocks for input
 - Thus, can merge at most $\lceil \frac{B-b}{b} \rceil$ sorted runs in each merge pass
- Analysis:**
 - N = number of pages in file to be sorted
 - B = number of available buffer pages
 - b = number of pages of each buffer block
 - N_b = number of initial sorted runs = $\lceil N/B \rceil$
 - F = number of runs that can be merged at each merge pass = $\lceil \frac{B}{b} \rceil - 1$
 - Number of passes = $\lceil \log_F(N_b) \rceil + 1$

- Set aside one buffer block for output of merge. ($B - b$). One buffer block per input run ($\frac{B-b}{b}$).

- Means can merge** at most floor($\frac{B-b}{b}$) sorted runs in each merge pass.

Blocked I/O: Sorting 36-page data with $B=6$



Overall:

- Larger buffer blocks, (lower average page I/O cost).
- But num. passes increase, num page I/O increase.

- Decrease per-page I/O cost tradeoffs Increase No. of page I/Os.**

Sorting using B^+ -trees

- When table to be sorted has existing B^+ -tree index on sorting attributes.
 - Format 1: Sequentially scan leaf pages of B^+ -tree.
 - Format 2/3: Sequentially scan leaf pages of B^+ -tree. For each leaf page visited, retrieve data records using RIDs.

Selection $\sigma_p(R)$

Select rows from Relation R that satisfy selection predicate p

- Index Matching:** Index **matches** selection predicate if index can be used to evaluate it. Consider Hash index, and B^+ Tree index.

Access Path

- Access path** refers to the given way of accessing data records/entries.
 - Table scan:** scan all data pages.
 - Index scan:** scan all index pages.
 - Index Intersection:** Combine results from multiple index scans (e.g. intersect, union).
- Index scan/intersect follow by **RID lookups** to retrieve data records.

Selectivity of Access Path

Selectivity of Access Path:

Number of index & data pages retrieved to access data record/entries

Most Selective Access Path:

(Smallest selectivity), retrieves fewest pages.

$Q_1: \text{select weight from } R \text{ where weight between 51 and 59}$
 $Q_2: \text{select * from } R \text{ where weight between 51 and 59}$

- Most Selective Access Path minimizes cost of data retrieval.**

Covering Index

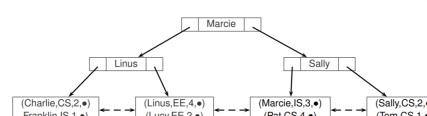
- An index I is a **covering index** for a query Q if all the attributes referenced in Q are part of the key or include column(s) of /
- Q can be evaluated using I without any RID lookup
- Such an evaluation plan is known as **index-only plan**

Example:

- Consider query Q :
 $\text{select major from Student where name = 'Bob'}$
 - An index on (name, major) is a covering index for Q
 - An index on (name) is not a covering index for Q
 - An index on (name) with include columns (address, major) is a covering index for Q

B^+ -Tree: Include Columns (In Index)

- Relation: Student(sid, name, major, year, address).
- B^+ -tree index on Student:
`create index stu_name_index on Student (name) include (major, year);`

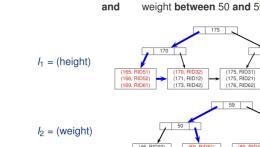


Query: `select major from Student where name = 'Lucy'`;

B^+ -Tree: Selection Evaluation

B^+ -tree : Index Intersection

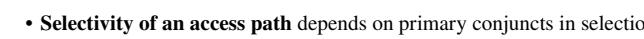
`select height, weight from Student where height between 164 and 170 and weight between 50 and 59`



$h = (\text{height})$
 $l_2 = (\text{weight})$



`select name from Student where weight between 55 and 59 and age > 20`



`Data Pages`

CNF Predicates

- Selectivity of an access path** depends on primary conjuncts in selection condition (w.r.t. index involved.)
- Each conjunct acts as filter on table, fraction of tuples satisfying given conjunct called reduction factor.

- A **term** is of the form $R.A \text{ op } c$ or $R.A_i \text{ op } R.A_j$

- A **conjunct** consists of one or more terms connected by \vee

- A conjunct that contains \vee is said to be **disjunctive** (or contains a disjunction)

- A **conjunctive normal form (CNF) predicate** consists of one or more conjuncts connected by \wedge



- Non-equality Comparison Operators:** $<$, \leq , $>$, \geq , $<>$, between, in.

Tree Index matching CNF Selection

- Determines if Index is useful / appropriate, if index can be used to retrieve just the tuples that satisfy the condition.

B^+ -tree : Matching predicates

- B^+ -tree index $I = (K_1, K_2, \dots, K_n)$
- Non-disjunctive CNF predicate p
- I matches p if p is of the form:

$$(K_1 = c_1) \wedge \dots \wedge (K_{i-1} = c_{i-1}) \wedge (K_i \text{ op } c_i), i \in [1, n]$$

zero or more equality predicates

where 1. (K_1, \dots, K_n) is a prefix of the key of I , and

2. there is at most one non-equality comparison operator which must be on the last attribute of the prefix (i.e., K_i)

Note: this definition is stronger than R&G's definition

B^+ -tree : Matching predicates (cont.)

Example: Which predicates does $I = (age, weight, height)$ match?

- age ≥ 20 Yes, match
- weight = 80 No
- (age = 20) \wedge (weight = 70) Yes, match
- (age = 20) \wedge (weight < 70) Yes, match
- (age > 20) \wedge (weight = 70) No
- (age = 20) \wedge (height = 170) Yes, match
- (height > 180) \wedge (weight = 65) \wedge (age = 20) Yes, match
- (age = 20) \wedge (weight ≤ 65) \wedge (height = 180) No

Hash Index matching CNF Selection

Hash Index: Matching predicates

Determines if entry lies in the bucket.

- Hash index $I = (K_1, K_2, \dots, K_n)$
- Non-disjunctive CNF predicate p
- I matches p if p is of the form:

$$(K_1 = c_1) \wedge \dots \wedge (K_n = c_n)$$

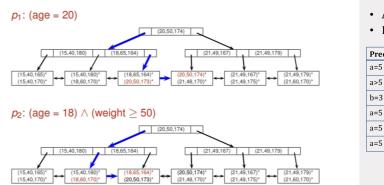
Hash Index: Matching predicates (cont.)

Example: Which predicates does $I = (age, weight, height)$ match?

- age ≥ 20 Yes, match
- weight = 80 No
- (age = 20) \wedge (weight < 70) Yes, match
- (age > 20) \wedge (weight = 70) No
- (age = 20) \wedge (height = 170) Yes, match
- (height > 180) \wedge (weight = 65) \wedge (age = 20) No
- (age = 20) \wedge (weight ≤ 65) \wedge (height = 180) Yes, match

Examples of Index matching CNF Selection

Example: B⁺-tree on (age,weight,height)



Primary and Covered Conjuncts

- Primary Conjuncts:** Subset of conjuncts in selection predicate p that index I matches.
- In general, only subset of conjuncts of predicate matches index.
- Covered Conjunct:** Conjunct C in predicate p covered if all attributes in C appear in the key, or *include column(s)* of index I .
- Primary conjuncts subset of covered conjuncts.

Cost of Evaluation of Selection Predicate p

Notation	Meaning
r	relational algebra expression
$ r $	number of tuples in output of r
$ r $	number of pages in output of r
b_d	number of data records that can fit on a page
b_i	number of data entries that can fit on a page
F	average fanout of B ⁺ -tree index (i.e., number of pointers to child nodes)
h	height of B ⁺ -tree index (i.e., number of levels of internal nodes)
$h = \lceil \log_F(\lceil \frac{ R }{b_i} \rceil) \rceil$	if format-2 index on table R
B	number of available buffer pages

Cost of B⁺-tree Index Evaluation of p

Let p' = primary conjuncts of p , p_c = covered conjuncts of p

1. Navigate internal nodes to locate first leaf page

$$Cost_{internal} = \begin{cases} \lceil \log_F(\lceil \frac{||R||}{b_i} \rceil) \rceil & \text{if } I \text{ is a format-1 index,} \\ \lceil \log_F(\lceil \frac{||R||}{b_i} \rceil) \rceil & \text{otherwise.} \end{cases}$$

2. Scan leaf pages to access all qualifying data entries

$$Cost_{leaf} = \begin{cases} \lceil \frac{||\sigma_{p'}(R)||}{b_d} \rceil & \text{if } I \text{ is a format-1 index,} \\ \lceil \frac{||\sigma_{p'}(R)||}{b_i} \rceil & \text{otherwise.} \end{cases}$$

3. Retrieve qualified data records via RID lookups

$$Cost_{rid} = \begin{cases} 0 & \text{if } I \text{ is covering or format-1 index,} \\ ||\sigma_{p_c}(R)|| & \text{otherwise.} \end{cases}$$

Cost of RID lookups could be reduced by first sorting the RIDs

$$\lceil \frac{||\sigma_{p_c}(R)||}{b_d} \rceil \leq Cost_{rid} \leq \min\{||\sigma_{p_c}(R)||, |R|\}$$

Cost of Hash Index Evaluation of p

- Let p' = primary conjuncts of p
- For format-1 index
 - Cost to retrieve data records: at least $\lceil \frac{||\sigma_{p'}(R)||}{b_d} \rceil$
- For format-2 index
 - Cost to retrieve data entries: at least $\lceil \frac{||\sigma_{p'}(R)||}{b_i} \rceil$
 - Cost to retrieve data records = $\begin{cases} 0 & \text{if } I \text{ is a covering index,} \\ \lceil \frac{||\sigma_{p_c}(R)||}{b_d} \rceil & \text{otherwise.} \end{cases}$

Evaluating Non-Disjunctive / Disjunctive Conjuncts

Evaluating Non-Disjunctive Conjuncts

- Consider the query $\sigma_p(R)$, where $p = (\text{age} = 21) \wedge (\text{weight} \geq 70) \wedge (\text{height} = 180)$
- Suppose the available unclustered indexes on R are
 - a hash index H_{age} on (age), and
 - a B⁺-tree index T_{weight} on (weight)
- What are the possible strategies to evaluate the following predicates?

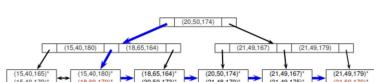
- $p_1 = (\text{age} = 21) \vee (\text{height} = 180)$
- $p_2 = ((\text{age} = 21) \vee (\text{height} = 180)) \wedge (\text{weight} \geq 70)$
- $p_3 = (\text{age} = 21) \vee (\text{weight} \geq 70)$

Possible strategies to evaluate Disjunctive / Non-Disjunctive predicates

- File scan, Use both (fetch RID, take Union), Use B+ tree etc.

Evaluating Disjunctive Conjuncts

- Suppose the available unclustered indexes are
 - a hash index H_{age} on (age), and
 - a B⁺-tree index T_{weight} on (weight)
- What are the possible strategies to evaluate the following predicates?



Example

- B⁺-tree index $I = (\text{age}, \text{weight}, \text{height})$, Format 2
- Query: `select * from R where p`
 - $p = (\text{age} \geq 18) \wedge (\text{weight} \geq 60) \wedge (\text{height} \geq 30)$
 - $p = (\text{age} \geq 18) \wedge (\text{height} \geq 30) \wedge (\text{weight} \geq 60)$
 - $p = (\text{age} \geq 18) \wedge (\text{weight} \geq 60)$
- $||R|| = 12$, $||\sigma_p(R)|| = 9$, $||\sigma_{p_c}(R)|| = 2$
- $b_d = 2$, $b_i = 2$, Height of $I = 2$
- Evaluation cost of p using $I = 2 + \lceil \frac{9}{2} \rceil + 2 = 9$